

netstat
network statistics

netstat

netstat (network statistics) is a command line tool for monitoring network connections both incoming and outgoing as well as viewing routing tables, interface statistics.

- Unix-like Operating Systems and also available on Windows OS as well
- Useful in terms of network troubleshooting and performance measurement.
- basic network service debugging tools
- Tells what ports are open and whether any programs are listening on ports.
- Very import tool for linux network administrators and system administrators
- troubleshoot their network related problems and determine network traffic performance.

netstat

Active Internet connections (w/o servers)

Proto	Recv-Q	Send-Q	Local Address	Foreign Address	State
tcp	0	0	prolin1:31027	prolin1:5500	TIME_WAIT
tcp	4	0	prolin1:1521	applin1:40205	ESTABLISHED
tcp	0	0	prolin1:1522	prolin1:39957	ESTABLISHED
tcp	0	0	prolin1:3938	prolin1:31017	TIME_WAIT
tcp	0	0	prolin1:1521	prolin1:21545	ESTABLISHED

- The above output goes on to show all the open sockets
- connection has to have a source and a destination, called *local* and *remote* address. The end points could be on the same server; or on different servers.
- the programs connect to the ~~same server~~. For instance, if two processes communicate among each other, the local and remote addresses will be the same, as you can see in the first line – the local and remote addresses are both the sever “prolin1”.
- However, the processes communicate over a *port*, which will be different. This port is shown next to the host name after the “:” (colon) mark
- The user program sends the data to be sent across the socket to a queue and the receiver reads from a queue at the remote end.

1. The leftmost column “ **Proto** ” shows the type of the connection – tcp in this case.
 2. The column **Recv-Q** shows the bytes of data in the queue to be sent to the user program that established the connection. This value should be as close to 0 as possible. In busy servers this value will be more than 0 but shouldn’t be very large. The **Send-Q** column denotes the bytes in the queue to be sent to the remote program, i.e. the remote program has acknowledged receiving it
 3. **Local Address** is source of the connection and the port number of the program.
-
4. **Foreign Address** is the destination host and port number. In the first line, both the source and destination are on the same host: proline 1. The connection is simply waiting. The second line shows an established connection between proln1 going to the port 40205 of the host appln1.
 5. The column **State** shows the status of the connection. Here are some common values.
 - **ESTABLISHED** – that the connection has been established. It does not mean that any data is flowing between the endpoints; merely that the end points have talked to each other.
 - **CLOSED** – the connection has been closed, i.e. not used now.
 - **TIME_WAIT** – the connection is being closed but there are still packets in the network that are being handled.
 - **CLOSE_WAIT** – the remote end has shutdown and has asked to close the connection.

The **-p** option shows the process information

```
netstat -p
```

Proto	Recv-Q	Send-Q	Local Address	Foreign Address	State	PID/Program name
tcp	0	0	prolin1:1521	prolin1:33303	ESTABLISHED	1327/oraclePROPRD1
tcp	0	0	prolin1:1521	applin1:51324	ESTABLISHED	13827/oraclePROPRD1
tcp	0	0	prolin1:1521	prolin1:33298	ESTABLISHED	32695/tnlsnr
tcp	0	0	prolin1:1521	prolin1:32544	ESTABLISHED	15251/oracle+ASM
tcp	0	0	prolin1:1521	prolin1:33331	ESTABLISHED	32695/tnlsnr

This clearly shows the process IP and the process name in the last column.

To find out the network statistics for various interfaces, use the -i option

```
netstat -i
```

Kernel Interface table

Iface	MTU	Met	RX-OK	RX-ERR	RX-DRP	RX-OVR	TX-OK	TX-ERR	TX-DRP	TX-OVR	Flg
eth0	1500	0	6860659	0	0	0	2055833	0	0	0	BMRU
eth8	1500	0	2345	0	0	0	833	0	0	0	BMRU
lo	16436	0	14449079	0	0	0	14449079	0	0	0	LRU

This shows the different interfaces present in the server (eth0, eth8, etc.) and the metrics associated with the interface.

- **RX-OK** shows the number of packets successfully received (for this interface)
- **RX-ERR** shows number of errors
- **RX-DRP** shows packets dropped
- **RX-OVR** shows packets overrun

The next sets of columns (TX-OK, TX-ERR, etc.) show the corresponding stats for send data.

Flg column is a composite value of the property of the interface. Each letter indicates a specific property. The following table provides an explanation of the letters.

B – Broadcast

M – Multicast

R – Running

U – Up

O – ARP Off

P – Point to Point Connection

L – Loopback

m – Master

You can use the `--interface` (note: there are *two* hyphens, not one) option to display the same for a specific interface.

```
netstat --interface=eth0
```

Kernel Interface table

Iface	MTU	Met	RX-OK	RX-ERR	RX-DRP	RX-OVR	TX-OK	TX-ERR	TX-DRP	TX-OVR	Flg
eth0	1500	0	277903459	0	0	0	170897632	0	0	0	BMsRU

```
netstat -i -e
```

Kernel Interface table

```
eth0    Link encap:Ethernet HWaddr 00:13:72:CC:EB:00
        inet addr:10.14.106.0 Bcast:10.14.107.255 Mask:255.255.252.0
        inet6 addr: fe80::213:72ff:fecc:eb00/64 Scope:Link
              UP BROADCAST RUNNING MULTICAST MTU:1500 Metric:1
              RX packets:6861068 errors:0 dropped:0 overruns:0 frame:0
              TX packets:2055956 errors:0 dropped:0 overruns:0 carrier:0
              collisions:0 txqueuelen:1000
              RX bytes:3574788558 (3.3 GiB) TX bytes:401608995 (383.0 MiB)
              Interrupt:169
```

Does the output seem familiar? It should; it's the same as the output of the **ifconfig**.

If you'd rather see the output showing IP addresses instead of host names, use the **-n** option.

The **-s** option shows the summary statistics of each protocol, rather than showing the details of each connection. This can be combined with the protocol specific flag. For instance **-u** shows the stats related to the UDP protocol.

The **-s** option shows the summary statistics of each protocol, rather than showing the details of each connection. This can be combined with the protocol specific flag. For instance **-u** shows the stats related to the UDP protocol.

netstat -s -u

Udp:

12764104 packets received

600849 packets to unknown port received.

0 packet receive errors

13455783 packets sent

Similarly, to see the stats for tcp, use **-t** and for raw, **-r**.

One of the really useful options is the display of the routing table, the **-r** option.

```
# netstat -r
```

Kernel IP routing table

Destination	Gateway	Genmask	Flags	MSS	Window	irtt	Iface
10.20.191.0	*	255.255.255.128	U	0	0	0	bond0
172.22.13.0	*	255.255.255.0	U	0	0	0	eth9
169.254.0.0	*	255.255.0.0	U	0	0	0	eth9
default	10.20.191.1	0.0.0.0	UG	0	0	0	bond0

The second column of netstat output— **Gateway**—shows the gateway to which the routing entry points. If no gateway is used, an asterisk is printed instead. The third column— **Genmask**—shows the “generality” of the route, i.e., the **network mask for this route**.

The fourth column, **Flags**, displays the following flags that describe the route:

- **G** means the route uses a gateway.
- **U** means the interface to be used is up (available).
- **H** means only a single host can be reached through the route. For example, this is the case for the loopback entry 127.0.0.1.
- **D** means this route is dynamically created.
- **!** means the route is a reject route and data will be dropped.

The next three columns show the **MSS**, **Window**, and **irtt** that will be applied to TCP connections established via this route.

- **MSS** stands for Maximum Segment Size – the size of the largest datagram for transmission via this route.
- **Window** is the maximum amount of data the system will accept in a single burst from a remote host for this route.
- **irtt** stands for Initial Round Trip Time.
- The TCP protocol has a built-in reliability check. If a data packet fails during transmission, it's re-transmitted. The protocol keeps track of how long it takes for the data to reach the destination and acknowledgement to be received. If the acknowledgement does not come within that timeframe, the packet is retransmitted. The amount of time the protocol has to wait before re-transmitting is set for the interface once (which can be changed) and that value is known as *initial round trip time*. A value of 0 means the default value is used.

Finally, the last field displays the network interface that this route will use.

Listing all the LISTENING Ports of TCP and UDP connections

```
# netstat -a | more
```

Active Internet connections (servers and established)

Proto	Recv-Q	Send-Q	Local Address	Foreign Address	State
tcp	0	0	*:sunrpc	*:*	LISTEN
tcp	0	52	192.168.0.2:ssh	192.168.0.1:egs	ESTABLISHED
tcp	1	0	192.168.0.2:59292	www.gov.com:http	CLOSE_WAIT
tcp	0	0	localhost:smtp	*:*	LISTEN
tcp	0	0	*:59482	*:*	LISTEN
udp	0	0	*:35036	*:*	
udp	0	0	*:npmp-local	*:*	

Active UNIX domain sockets (servers and established)

Proto	RefCnt	Flags	Type	State	I-Node	Path
unix	2	[ACC]	STREAM	LISTENING	16972	/tmp/orbit-root/linc-76b-0-6fa08790551
unix	2	[ACC]	STREAM	LISTENING	17149	/tmp/orbit-root/linc-794-0-7058d58416
unix	2	[ACC]	STREAM	LISTENING	17161	/tmp/orbit-root/linc-792-0-546fe905321
unix	2	[ACC]	STREAM	LISTENING	15938	/tmp/orbit-root/linc-74b-0-415135cb6ae

Listing TCP Ports connections

Listing only TCP (Transmission Control Protocol) port connections using netstat -at.

```
# netstat -at

Active Internet connections (servers and established)
Proto Recv-Q Send-Q Local Address          Foreign Address        State
tcp      0      0  *:ssh                  *:*                  LISTEN
tcp      0      0  localhost:ipp           *:*                  LISTEN
tcp      0      0  localhost:smtp          *:*                  LISTEN
tcp      0      52 192.168.0.2:ssh        192.168.0.1:egs      ESTABLISHED
tcp      1      0  192.168.0.2:59292     www.gov.com:http     CLOSE_WAIT
```

Listing UDP Ports connections

Listing only UDP (User Datagram Protocol) port connections using netstat -au.

```
# netstat -au

Active Internet connections (servers and established)
Proto Recv-Q Send-Q Local Address          Foreign Address        State
udp      0      0 *:35036                *:*
udp      0      0 *:npmp-local            *:*
udp      0      0 *:mdns                *:*
```

Listing all LISTENING Connections

Listing all active listening ports connections with netstat -l.

```
# netstat -l

Active Internet connections (only servers)
Proto Recv-Q Send-Q Local Address           Foreign Address         State
tcp      0      0 *:sunrpc                 *:*                  LISTEN
tcp      0      0 *:58642                  *:*                  LISTEN
tcp      0      0 *:ssh                    *:*                  LISTEN
udp      0      0 *:35036                  *:* 
udp      0      0 *:npnmp-local            *:* 

Active UNIX domain sockets (only servers)
Proto RefCnt Flags       Type      State         I-Node  Path
unix    2      [ ACC ]     STREAM   LISTENING   16972  /tmp/orbit-root/linc-76b-0-6fa08790553
unix    2      [ ACC ]     STREAM   LISTENING   17149  /tmp/orbit-root/linc-794-0-7058d584166
unix    2      [ ACC ]     STREAM   LISTENING   17161  /tmp/orbit-root/linc-792-0-546fe905321
unix    2      [ ACC ]     STREAM   LISTENING   15938  /tmp/orbit-root/linc-74b-0-415135cb6ae
```

Listing all TCP Listening Ports

Listing all active listening TCP ports by using option netstat -lt.

```
# netstat -lt

Active Internet connections (only servers)
Proto Recv-Q Send-Q Local Address          Foreign Address        State
tcp      0      0  *:dctcp                *:*                  LISTEN
tcp      0      0  *:mysql                *:*                  LISTEN
tcp      0      0  *:sunrpc               *:*                  LISTEN
tcp      0      0  *:munin                *:*                  LISTEN
tcp      0      0  *:ftp                  *:*                  LISTEN
tcp      0      0  localhost.localdomain:ipp  *:*                  LISTEN
tcp      0      0  localhost.localdomain:smtp *:*                  LISTEN
tcp      0      0  *:http                 *:*                  LISTEN
tcp      0      0  *:ssh                  *:*                  LISTEN
tcp      0      0  *:https               *:*                  LISTEN
```

Listing all UDP Listening Ports

Listing all active listening UDP ports by using option netstat -lu.

```
# netstat -lu

Active Internet connections (only servers)
Proto Recv-Q Send-Q Local Address          Foreign Address        State
udp      0      0  *:39578                *:*
udp      0      0  *:meregister          *:*
udp      0      0  *:vpps-qua            *:*
udp      0      0  *:openvpn             *:*
udp      0      0  *:mdns                *:*
udp      0      0  *:sunrpc              *:*
udp      0      0  *:ipp                 *:*
udp      0      0  *:60222              *:*
udp      0      0  *:mdns                *:*
```

Showing Statistics by Protocol

Displays statistics by protocol. By default, statistics are shown for the TCP, UDP, ICMP, and IP protocols. **The -s parameter can be used to specify a set of protocols.**

Showing statistics of only **TCP protocol** by using option **netstat -st**.

Showing statistics of only **UDP Protocol** **netstat -su**

```
# netstat -s

Ip:
    2461 total packets received
    0 forwarded
    0 incoming packets discarded
    2431 incoming packets delivered
    2049 requests sent out

Icmp:
    0 ICMP messages received
    0 input ICMP message failed.
    ICMP input histogram:
        1 ICMP messages sent
        0 ICMP messages failed
    ICMP output histogram:
        destination unreachable: 1

Tcp:
    159 active connections openings
    1 passive connection openings
    4 failed connection attempts
    0 connection resets received
    1 connections established
    2191 segments received
    1745 segments send out
    24 segments retransmited
    0 bad segments received.
    4 resets sent

Udp:
    243 packets received
    1 packets to unknown port received.
    0 packet receive errors
```

Displaying IPv4 and IPv6 Information

```
# netstat -g

IPv6/IPv4 Group Memberships
Interface      RefCnt Group
-----
lo            1      all-systems.mcast.net
eth0          1      224.0.0.251
eth0          1      all-systems.mcast.net
lo            1      ff02::1
eth0          1      ff02::202
eth0          1      ff02::1:ffb4:da21
eth0          1      ff02::1
```

Displaying RAW Network Statistics

```
# netstat --statistics --raw

Ip:
    62175683 total packets received
        52970 with invalid addresses
        0 forwarded

Icmp:
    875519 ICMP messages received
        destination unreachable: 901671
        echo request: 8
        echo replies: 16253

IcmpMsg:
        InType0: 83

IpExt:
        InMcastPkts: 117
```

Ifconfig (Interface configuration)

ifconfig

The ifconfig command shows the details of the network interface(s) defined in the system.

configure, manage and query network interface parameters via command line

The most common option is **-a** , which shows all the interfaces.

ifconfig -a

-a display all interfaces which are currently available, even if down

The usual name of the primary Ethernet network interface is eth0. To find out the details of a specific interface, e.g. eth0, you can use:

ifconfig eth0

The output is show below

```
# ifconfig eth0
```

```
eth0      Link encap:Ethernet  HWaddr 00:13:72:CC:EB:00
          inet  addr:10.14.106.0  Bcast:10.14.107.255  Mask:255.255.252.0
          inet6 addr: fe80::213:72ff:fecc:eb00/64 Scope:Link
                     UP BROADCAST RUNNING MULTICAST  MTU:1500  Metric:1
                     RX packets:6853429 errors:0 dropped:0 overruns:0 frame:0
                     TX packets:2055296 errors:0 dropped:0 overruns:0 carrier:0
                     collisions:0 txqueuelen:1000
                     RX bytes:3573539098 (3.3 GiB)  TX bytes:401474390 (382.8 MiB)
                     Interrupt:169
```

Type of Media

IP

Address

Metrics on Transmission and Reception of Data

MAC

Netmask

Here are some key parts of the output:

- **Link encap:** the type of the hardware physical medium supported by this interface (Ethernet, in this case)
- **HWaddr:** the unique identifier of the NIC card. Every NIC card has a unique identifier assigned by the manufacturer, called MAC or MAC address. The IP address is attached to the MAC to the server. If this IP address is changed, or this card is moved from this server to a different one, the MAC does not change.
- **Mask:** the netmask
- **inet addr:** the IP address attached to the interface
- **RX packets:** the number of packets received by this interface
- **TX packets:** the number of packets sent

The command is not just used to check the settings; it's used to configure and manage the interface as well. Here is a short list of parameters and options for this command:

up/down – enables or disables a specific interface. You can use the down parameter to shutdown an interface (or disable it):

```
# ifconfig eth0 down
```

Similarly to bring it up (or enable) it, you would use:

```
# ifconfig eth0 up
```

add – sets a specific IP address for the interface. To set an IP address of 192.168.1.101 to the interface eth0, you would issue:

```
# ifconfig eth0 add 192.168.1.101
```

netmask – sets the netmask parameter of the interface. Here is an example where you can set the netmask of the eth0 interface to 255.255.255.0

```
# ifconfig eth0 netmask 255.255.255.0
```

How to Assign a Broadcast to Network Interface

Using the “broadcast” argument with an interface name will set the broadcast address for the given interface. For example, “ifconfig eth0 broadcast 172.16.25.63” command sets the broadcast address to an interface eth0.

ifconfig eth0 broadcast 172.16.25.63

How to Assign a IP, Netmask and Broadcast to Network Interface

ifconfig eth0 172.16.25.125 netmask 255.255.255.244 broadcast 172.16.25.63

How to Change MTU for an Network Interface

The “mtu” argument set the maximum transmission unit to an interface. The MTU allows you to set the limit size of packets that are transmitted on an interface.

ifconfig eth0 mtu 1000

How to Enable Promiscuous Mode

ifconfig eth0 promisc

How to Disable Promiscuous Mode

ifconfig eth0 -promisc

How to Add New Alias to Network Interface

ifconfig eth0:0 172.16.25.127

Next, verify the newly created alias network interface address, by using “**ifconfig eth0:0**” command.

```
eth0:0      Link encap:Ethernet  HWaddr 00:01:6C:99:14:68
            inet addr:172.16.25.123  Bcast:172.16.25.63  Mask:255.255.255.240
                  UP BROADCAST RUNNING MULTICAST  MTU:1500  Metric:1
                  Interrupt:17
```

How to Remove Alias to Network Interface

ifconfig eth0:0 down

How to Change the MAC address of Network Interface

To change the MAC (Media Access Control) address of an eth0 network interface, use the following command with argument “hw ether”

ifconfig eth0 hw ether AA:BB:CC:DD:EE:FF

IP Command

IP Command

The ip command is used to assign an address to a network interface and/or configure network interface parameters on Linux operating systems. This command replaces old good and now deprecated ifconfig command on modern Linux distributions.

Use this command to display and **configure the network parameters** for host interfaces.

Displays info about all network interfaces

Type the following command to list and show all ip address associated on on all network interfaces:

```
ip a
```

OR

```
ip addr
```

```
user1@vm1:~$ ip a
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 16436 qdisc noqueue state UNKNOWN
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        inet6 ::1/128 scope host
            valid_lft forever preferred_lft forever
2: eth0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc pfifo_fast state UP qlen 1000
    link/ether 08:00:27:d4:45:68 brd ff:ff:ff:ff:ff:ff
    inet 10.0.2.15/24 brd 10.0.2.255 scope global eth0
        inet6 fe80::a00:27ff:fed4:4568/64 scope link
            valid_lft forever preferred_lft forever
```

You can select between IPv4 and IPv6 using the following syntax:

Only show TCP/IP IPv4

ip -4 a

Only show TCP/IP IPv6

ip -6 a

It is also possible to specify and list particular interface TCP/IP details:

Only show eth0 interface

ip a show eth0

ip a list eth0

ip a show dev eth0

Only show running interfaces

ip link ls up

Assigns the IP address to the interface

The syntax is as follows to add an IPv4/IPv6 address:

```
ip a add {ip_addr/mask} dev {interface}
```

To assign 192.168.1.200/255.255.255.0 to eth0, enter:

```
ip a add 192.168.1.200/255.255.255.0 dev eth0
```

OR

```
ip a add 192.168.1.200/24 dev eth0
```

Adding the broadcast address on the interface

By default, the ip command does not set any broadcast address unless explicitly requested. So syntax is as follows to set broadcast ADDRESS:

```
ip addr add brd {ADDRESS-HERE} dev {interface}
```

```
ip addr add broadcast {ADDRESS-HERE} dev {interface}
```

```
ip addr add broadcast 172.20.10.255 dev dummy0
```

Remove / Delete the IP address from the interface

The syntax is as follows to remove an IPv4/IPv6 address:

```
ip a del {ipv6_addr_OR_ipv4_addr} dev {interface}
```

To delete 192.168.1.200/24 from eth0, enter:

```
ip a del 192.168.1.200/24 dev eth0
```

Flush the IP address from the interface

You can delete or remove an IPv4/IPv6 address one-by-one as described above. However, the flush command can remove as flush the IP address as per given condition. For example, you can delete all the IP addresses from the private network 192.168.2.0/24 using the following command:

```
ip -s -s a f to 192.168.2.0/24
```

Sample outputs:

```
2: eth0      inet 192.168.2.201/24 scope global secondary eth0
2: eth0      inet 192.168.2.200/24 scope global eth0
*** Round 1, deleting 2 addresses ***
*** Flush is complete after 1 round ***
```

You can disable IP address on all the ppp (Point-to-Point) interfaces:

```
ip -4 addr flush label "ppp*"
```

change the state of the device to UP or DOWN?

The syntax is as follows:

```
ip link set dev {DEVICE} {up|down}
```

To make the state of the device eth1 down, enter:

```
ip link set dev eth1 down
```

To make the state of the device eth1 up, enter:

```
ip link set dev eth1 up
```

change the MTU of the device

```
ip link set mtu {NUMBER} dev {DEVICE}
```

To change the MTU of the device eth0 to 9000, enter:

```
ip link set mtu 9000 dev eth0
```

ip route: Routing table management commands

Use the following command to manage or manipulate the kernel routing table.

Show routing table

To display the contents of the routing tables:

```
ip r
```

```
ip r list
```

```
ip r list [options]
```

```
ip route
```

Sample output

```
192.168.1.0/24 dev eth1  proto kernel  scope link  src 192.168.1.10
```

Display routing for 192.168.1.0/24:

```
ip r list 192.168.1.0/24
```

Sample outputs:

```
192.168.1.0/24 dev eth1  proto kernel  scope link  src 192.168.1.10
```

Add a new route

The syntax is:

```
ip route add {NETWORK/MASK} via {GATEWAYIP}
```

```
ip route add {NETWORK/MASK} dev {DEVICE}
```

```
ip route add default {NETWORK/MASK} dev {DEVICE}
```

```
ip route add default {NETWORK/MASK} via {GATEWAYIP}
```

Add a plain route to network 192.168.1.0/24 via gateway 192.168.1.254:

```
ip route add 192.168.1.0/24 via 192.168.1.254
```

Delete a route

The syntax is as follows to delete default gateway:

```
ip route del default
```

```
ip route del 192.168.1.0/24 dev eth0
```

Ping

Ping

Test the connection between the local server/computer and a remote UNIX server. The ping command sends ICMP Echo Request (ECHO_REQUEST) packets to the host once per second. Each packet that is echoed back via an ICMP Echo Response packet is written to the standard output, including round-trip time. Use this command to find out:

1. If a remote server is up and running.
2. Test network problems.
3. Test network connectivity to your local or remote gateways.
4. Verify network problems.

ping google.com

```
user1@vm1:~$ ping google.com
PING google.com (74.125.28.102) 56(84) bytes of data.
64 bytes from pc-in-f102.1e100.net (74.125.28.102): icmp_req=1 ttl=39 time=246 m
s
64 bytes from pc-in-f102.1e100.net (74.125.28.102): icmp_req=2 ttl=39 time=276 m
s
64 bytes from pc-in-f102.1e100.net (74.125.28.102): icmp_req=3 ttl=39 time=311 m
s
64 bytes from pc-in-f102.1e100.net (74.125.28.102): icmp_req=4 ttl=39 time=293 m
s
64 bytes from pc-in-f102.1e100.net (74.125.28.102): icmp_req=5 ttl=39 time=448 m
s
64 bytes from pc-in-f102.1e100.net (74.125.28.102): icmp_req=6 ttl=39 time=298 m
s
```

time=69ms" This is the round trip time: the time between sending "Are you there?" and receiving "Yes I am!".

TTL : (Time to live)

A timer value included in packets sent over TCP/IP-based networks that tells the recipients how long to hold or use the packet or any of its included data before expiring and discarding the packet or data.

mdev

It's the standard deviation, essentially an average of how far each ping RTT is from the mean RTT. The higher **mdev** is, the more variable the RTT is (over time).

lcmp_seq :sequence number

lcmp_req:request number

Send limited number of pings

You can stop after sending 4 ECHO_REQUEST packets with the -c option as follows:

```
ping -c -4 yahoo.com
```

-c count -stop after sending count echo request packet, ping waits for count echo reply packets.

Avoid dns lookup

Pass the -n option to avoid to lookup symbolic names for host addresses i.e. numeric output only:

```
ping -n -c 4 google.com
```

```
user1@vm1:~$ ping -n -c 4 google.com
PING google.com (74.125.28.139) 56(84) bytes of data.
64 bytes from 74.125.28.139: icmp_req=1 ttl=39 time=289 ms
64 bytes from 74.125.28.139: icmp_req=2 ttl=39 time=296 ms
64 bytes from 74.125.28.139: icmp_req=3 ttl=39 time=294 ms
64 bytes from 74.125.28.139: icmp_req=4 ttl=39 time=309 ms
```

```
user1@vm1:~$ ping -c 4 google.com
PING google.com (74.125.28.101) 56(84) bytes of data.
64 bytes from pc-in-f101.1e100.net (74.125.28.101): icmp_req=1 ttl=39 time=287 ms
64 bytes from pc-in-f101.1e100.net (74.125.28.101): icmp_req=2 ttl=39 time=268 ms
64 bytes from pc-in-f101.1e100.net (74.125.28.101): icmp_req=3 ttl=39 time=297 ms
^C64 bytes from pc-in-f101.1e100.net (74.125.28.101): icmp_req=4 ttl=39 time=269 ms
```

Increase or Decrease the Time Interval Between Packets

default ping waits for 1 second before sending the next packet.

Increase Ping Time Interval

Example: Wait for 5 seconds before sending the next packet.

```
$ ping -i 5 IP
```

Decrease Ping Time Interval

Example: Wait 0.1 seconds before sending the next packet.

```
ping -i 0.1 IP
```

Note: Only super user can specify interval less than 0.2 seconds.

Check whether the local network interface is up and running

```
ping 0
```

```
ping localhost
```

```
ping 127.0.0.1
```

Flood the network

Super users can send hundred or more packets per second using -f option.

```
ping -f localhost
```

```
PING localhost (127.0.0.1) 56(84) bytes of data.
```

```
.^C
```

```
--- localhost ping statistics ---
```

```
427412 packets transmitted, 427412 received, 0% packet loss, time 10941ms
```

Audible ping

```
ping -a IP
```

Give beep when the peer is reachable

Find out the IP address

```
ping -c 1 google.com
```

Specify path for ping to send the packet

```
ping hop1 hop2 hop3 .. hopN destination
```

```
ping 192.168.3.33 192.168.7.1 192.168.4.45
```

Record and print route of how ECHO_REQUEST sent and ECHO_REPLY received

It records, and prints the network route through which the packet is sent and received.

```
ping -R 192.168.1.63
```

```
PING 192.168.1.63 (192.168.1.63) 56(84) bytes of data.
```

```
64 bytes from 192.168.1.63: icmp_seq=1 ttl=61 time=2.05 ms
```

```
RR: 192.168.9.118
```

```
192.168.3.25
```

```
192.168.10.35
```

```
192.168.1.26
```

```
192.168.1.63
```

```
192.168.1.63
```

```
192.168.10.4
```

```
192.168.3.10
```

```
192.168.4.25
```

Traceroute




traceroute utility prints out complete route to a particular destination.

Traceroute utility uses the TTL field in the IP header to achieve its operation.

TTL field describes how much hops a particular packet will take while traveling on network.

this effectively outlines the lifetime of the packet on network. This field is usually set to 32 or 64.

- Each time the packet is held on an intermediate router.
- decreases the TTL value by 1
- When a router finds the TTL value of 1 in a received packet then that packet is not forwarded but instead discarded.
- After discarding the packet, router sends an ICMP error message of “Time exceeded” back to the source from where packet generated.
- The ICMP packet that is sent back contains the IP address of the router.
- So now it can be easily understood that traceroute operates by sending packets with TTL value starting from 1 and then incrementing by one each time
- So traceroute incrementally fetches the IP of all the routers between the source and the destination.



How to run traceroute?

traceroute google.com

```
user1@vm1:~$ traceroute google.com
traceroute to google.com (74.125.28.101), 30 hops max, 60 byte packets
 1  10.0.2.2 (10.0.2.2)  0.440 ms  0.267 ms  0.267 ms
 2  * * *
 3  * * *
 4  * * *
 5  * * *
 6  * * *
 7  *^C
```

How To Read Traceroute's Output

The first line tells us the conditions that traceroute is operating under:

traceroute to google.com (173.194.38.137), 30 hops max, 60 byte packets

- It gives the specified host
- IP address that DNS returns for that domain
- maximum number of hops to check
- size of the packet that will be used.

- Each line gives the details of interaction with each router encountered.
- traceroute not only gives the IP addresses of the intermediate routers but also three round trip times for that particular router as for each router the traceroute command fires three packets.
- The “*” field in output
- This depicts that the required field could not be fetched.

- The maximum number of hops can be adjusted with the -m flag.
- host you are trying to route to is over 30 hops away.
- The maximum value you can set is 255.

```
traceroute -m 255 google.com
```

Adjust the size of the packet that is sent to each hop

```
traceroute google.com 70
```

traceroute to google.com (173.194.38.128), 30 hops max, **70 byte packets**

```
1 192.241.160.254 (192.241.160.254) 0.364 ms 0.330 ms 0.319 ms
2 192.241.164.237 (192.241.164.237) 0.284 ms 0.343 ms 0.321 ms
```

After the first line, each subsequent line represents a "**hop**", or **intermediate host** that your traffic must pass through to reach the computer represented by the host you specified.

Each line has the following format:

hop_number host_name (IP_address) packet_round_trip_times

hop_number: A sequential count of the number of degrees of separation the host is from your computer.

IP_address: This field contains the IP address for this network hop.

Roundtrip_times: The remainder of the line gives the round-trip times for a packet to the host and back again

3 nyk-b6-link.telia.net (62.115.35.101) 0.311 ms 0.302 ms 0.293 ms

By default, three packets are sent to each host

Configure the TTL value to start with

Traceroute utility is **flexible enough to accept the TTL value** that the user wants to start the utility with. By **default its value is 1** which means it starts off with the first router in the path but using the '**-f option**' (which expects the new value of TTL) a new value of the TTL field can be set.

```
traceroute google.com
```

```
traceroute to google.com (74.125.236.132), 30 hops max, 60 byte packets
```

```
1 220.224.141.129 (220.224.141.129) 89.181 ms 101.540 ms 101.503 ms
```

```
2 115.255.239.65 (115.255.239.65) 101.468 ms 101.431 ms 101.324 ms
```

```
3 124.124.251.245 (124.124.251.245) 121.373 ms 121.350 ms 158.694 ms
```

```
4 115.255.239.45 (115.255.239.45) 101.223 ms 141.135 ms 123.932 ms
```

```
5 72.14.212.118 (72.14.212.118) 123.867 ms 123.832 ms 123.802 ms
```

```
6 72.14.232.202 (72.14.232.202) 123.773 ms 123.742 ms 587.812 ms
```

```
7 216.239.48.179 (216.239.48.179) 587.723 ms 587.681 ms 587.642 ms
```

```
8 bom03s02-in-f4.1e100.net (74.125.236.132) 577.548 ms 577.524 ms 587.512 ms
```

```
$ traceroute google.com -f 8
```

```
traceroute to google.com (74.125.236.129), 30 hops max, 60 byte packets
```

```
8 bom03s02-in-f1.1e100.net (74.125.236.129) 96.961 ms 96.886 ms 96.849 ms
```

Configure Number of Queries per Hop

- the traceroute utility sends 3 packets per hop to provide 3 round trip times
- This default value of 3 is configurable using the option ‘-q’.

```
$ traceroute google.com -q 5
```

```
traceroute to google.com (173.194.36.46), 30 hops max, 60 byte packets
```

```
1  220.224.141.129 (220.224.141.129)  91.579 ms  91.497 ms  91.458 ms  91.422 ms  91.385 ms
```

```
2  115.255.239.65 (115.255.239.65)  91.356 ms  91.325 ms  98.868 ms  98.848 ms  98.829 ms
```

Configure Response Wait Time

- The time for which traceroute utility waits after issuing a probe can also be configured.
- The **-w** option expects a value which the utility will take as the response time to wait for.
- the wait time is 0.1 seconds and the traceroute utility was unable to wait for any response and it printed all the *'s.

```
traceroute google.com -w 0.1
```

```
traceroute to google.com (74.125.236.101), 30 hops max, 60 byte packets
```

```
1 * * *
```

```
2 * * *
```

```
3 * * *
```

```
..
```

```
26 * * *
```

```
27 * * *
```

```
28 * * *
```

```
29 * * *
```

```
30 * * *
```

Disable IP address and host name mapping

Traceroute provides an option through which the **mapping of IP addresses with host name** (that traceroute tries) **is disabled**. The option for doing this is '**-n**' . The following example illustrates this :

```
traceroute google.com -n

traceroute to google.com (173.194.36.7), 30 hops max, 60 byte packets

1  220.224.141.129  109.352 ms  109.280 ms  109.248 ms

2  115.255.239.65  131.633 ms  131.598 ms  131.573 ms

3  124.124.251.245  131.554 ms  131.529 ms  131.502 ms

4  115.255.239.45  131.478 ms  131.464 ms  199.741 ms
```

mtr(My Traceroute)

- It is a **powerful network diagnostic tool**
- combines the power of both **Ping** and **Traceroute** commands.
- diagnose and isolate network errors
- provide helpful **network status reports**.
- MTR works by sending ICMP packets by incrementally increasing the TTL value to find the route between the source and the given destination.
- MTR collects additional information regarding the **state, connection, and responsiveness of the intermediate hosts**.

Installation

On debian or Ubuntu systems use the following command:

```
$ sudo apt-get install mtr
```

On centos and fedora systems execute the following command:

```
$ yum install mtr
```

Execute mtr for a Domain

MTR works in two modes, a graphical mode (X11) and text based mode (ncurses). By default, mtr command runs in X11 mode.

```
mtr google.com
```

The above command will open up a GUI window, and display the results.

Launch Text Mode using –curses

Use the –curses option to run mtr in terminal mode.

```
mtr --curses google.com
```

Or

```
mtr -t google.com
```

packets **pings**

Host	Loss%	Snt	Last	Avg	Best	Wrst	StDev
1. mblaze.hilink	0.0%	20	3.2	2.7	2.1	3.8	0.5
2. 10.228.129.9	0.0%	20	187.7	61.8	41.6	187.7	31.3
3. 10.228.149.14	0.0%	20	112.1	60.2	33.2	112.1	17.8
4. 116.202.226.145	0.0%	20	57.9	63.2	35.2	147.6	24.4
5. 116.202.226.21	0.0%	20	35.4	70.4	35.4	211.8	48.6
6. 72.14.219.94	0.0%	20	58.9	74.6	43.4	231.2	44.2
7. 72.14.233.204	0.0%	20	46.9	69.8	40.3	222.5	41.9
8. 72.14.239.20	0.0%	20	94.1	259.2	68.8	3436.	748.2
9. 209.85.244.111	0.0%	20	86.4	97.5	72.1	232.2	34.3
10. google.com	0.0%	19	387.9	132.5	71.8	387.9	84.9

The above will run continuously in interactive mode.

Generate MTR reports using the following syntax:

`mtr -rw google.com`

HOST		Loss%	Snt	Last	Avg	Best	Wrst	StDev	
1. --	172.17.2.1	0.0%	10	1.0	1.0	0.9	1.1	0.0	
2. --	static-29.134.16.61-tataidc.co.in	0.0%	10	5.6	7.2	4.8	23.3	5.6	
3. --	???	100.0	10	0.0	0.0	0.0	0.0	0.0	
4. --	???	100.0	10	0.0	0.0	0.0	0.0	0.0	
5. --	10.117.227.50	0.0%	10	29.7	30.4	29.2	31.8	0.5	
6. --	static-98.110.93.111-tataidc.co.in	0.0%	10	124.2	84.4	34.9	124.2	33.5	
7. --	14.141.20.165.static-vsnl.net.in	0.0%	10	105.8	83.7	33.2	124.9	29.6	
8. --	???	100.0	10	0.0	0.0	0.0	0.0	0.0	
9. --	???	100.0	10	0.0	0.0	0.0	0.0	0.0	
10. --	115.114.142.137.static-Chennai.vsnl.net.in	10.0%		10	121.0	100.3	78.7	125.7	17.1
11. --	121.240.1.50	20.0%		10	116.5	93.1	67.9	116.5	18.4
12. --	72.14.233.204	10.0%		10	84.9	90.0	54.1	123.5	25.8
13. --	209.85.255.43	0.0%	9	57.2	66.9	28.7	119.3	31.8	
14. --	maa03s18-in-f46.1e100.net	0.0%	9	50.3	71.0	29.9	104.7	28.8	

The `r` option flag generates the report (short for `--report`)

The `w` option flag uses the long-version of the hostname (short for `--report-wide`).

```
mtr --report google.com (with out using w flag)
```

Start: Thu Apr 7 10:35:05 2016

HOST: HIFX-33 Loss% Snt Last Avg Best Wrst StDev

1.|-- 172.17.2.1 0.0% 10 1.1 1.0 0.9 1.1 0.0

2.1-- static-29.134.16.61-tatai 0.0% 10 6.2 9.9 5.2 44.7 12.2

3.|-- ??? 100.010 0.0 0.0 0.0 0.0 0.0

4.|-- ??? 100.010 0.0 0.0 0.0 0.0 0.0

5.|-- 10.117.227.50 0.0% 10 31.9 30.8 30.5 31.9 0.3

6.1-- static-98.110.93.111-tata 0.0% 10 114.0 94.4 57.4 116.6 17.4

7.1--14.141.20.165.static-vsni 0.0% 10.106.3.99.8.46.5.149.2.26.8

8 |-- ??? 100 010 00 00 00 00 00

9 |-- ??? 100 010 00 00 00 00 00

10 115 114 142 137 static-Ch 20 0% 10 112 6 113 2 105 8 142 4 12 2

11.1–121.240.1.50 20.0% 10.122.5.116.0.103.9.134.4.9.2

12 1-- 72 11 233 204 20.0% 10 113 8 111 1 89 1 119 1 7 7

13 209 85 255 43 10.0% 10 113 4 115 8 107 2 122 7 1 8

14.1 maa03c18_in_f14_1e100.net 20.0% 10.111.7.111.5.101.5.118.3.1.0

- report option which **sends 10 packets** to the host google.com and generates the output.
- Without the --report option, mtr will run continuously in an interactive environment.

The names for the hosts are determined by reverse DNS lookups

mtr --no-dns --report google.com

HOST: deleuze	Loss%	Snt	Last	Avg	Best	Wrst	StDev
1. 192.168.1.1	0.0%	10	2.2	2.2	2.0	2.7	0.2
2. 68.85.118.13	0.0%	10	8.6	11.0	8.4	17.8	3.0
3. 68.86.210.126	0.0%	10	9.1	12.1	8.5	24.3	5.2
4. 68.86.208.22	0.0%	10	12.2	15.1	11.7	23.4	4.4
5. 68.85.192.86	0.0%	10	17.2	14.8	13.2	17.2	1.3
6. 68.86.90.25	0.0%	10	14.2	16.4	14.2	20.3	1.9

MTR provides valuable statistics regarding the durability of that connection in the seven columns that follow.

The **Loss%** column shows the percentage of packet loss at each hop.

The **Snt** column counts the number of packets sent.

The next four columns **Last, Avg, Best, and Wrst** are all measurements of **latency in milliseconds**

- **Last** is the latency of the last packet sent.
- **Avg** is average latency of all packets.
- **Best** and **Wrst** display the best (shortest) and worst (longest) round trip time for a packet to this host.
- **StDev**, provides the standard deviation of the latencies to each host.
- The higher the standard deviation, the greater the difference is between measurements of latency

Verify Packet Loss

Rate Limit--common practice among the service provider to “Rate Limit” the ICMP traffic.

This can provide an illusion of packet loss

fact there is no loss.

to verify whether the loss is real or due to rate limiting, check the “Loss%” of the next hop.

If it shows 0.0%, then you can be sure that the “Loss%” reported is due to the ICMP rate limiting and not actual loss.

```
root@localhost:~# mtr --report www.google.com
```

HOST: example	Loss%	Snt	Last	Avg	Best	Wrst	StDev
1. 63.247.74.43	0.0%	10	0.3	0.6	0.3	1.2	0.3
2. 63.247.64.157	50.0%	10	0.4	1.0	0.4	6.1	1.8
3. 209.51.130.213	0.0%	10	0.8	2.7	0.8	19.0	5.7
4. aix.pr1.atl.google.com	0.0%	10	6.7	6.8	6.7	6.9	0.1
5. 72.14.233.56	0.0%	10	7.2	8.3	7.1	16.4	2.9
6. 209.85.254.247	0.0%	10	39.1	39.4	39.1	39.7	0.2
7. 64.233.174.46	0.0%	10	39.6	40.4	39.4	46.9	2.3
8. gw-in-f147.1e100.net	0.0%	10	39.6	40.5	39.5	46.7	2.2

loss continues for more than 1 hop, then it is possible that there is some packet loss.

Rate limiting and Packet Loss can happen concurrently

In that case take the lowest Loss% in a sequence as actual loss.

Improper Destination Host Networking

```
root@localhost:~# mtr --report www.google.com
HOST: localhost          Loss% Snt Last Avg Best Wrst StDev
 1. 63.247.74.43          0.0% 10  0.3  0.6  0.3  1.2  0.3
 2. 63.247.64.157          0.0% 10  0.4  1.0  0.4  6.1  1.8
 3. 209.51.130.213          0.0% 10  0.8  2.7  0.8 19.0  5.7
 4. aix.pr1.atl.google.com    0.0% 10  6.7  6.8  6.7  6.9  0.1
 5. 72.14.233.56          0.0% 10  7.2  8.3  7.1 16.4  2.9
 6. 209.85.254.247          0.0% 10 39.1 39.4 39.1 39.7  0.2
 7. 64.233.174.46          0.0% 10 39.6 40.4 39.4 46.9  2.3
 8. gw-in-f147.1e100.net    100.0 10  0.0  0.0  0.0  0.0  0.0
```

From the above example, it may look like the packets didn't reach the destination.
But it does reach the destination.
This may be a result of improperly configured host or firewall rules to drop ICMP packets.

Timeout and Return Route Issue

Sometimes, routers will discard the ICMP and it will be shown as ??? on the output.

Alternatively there can be a problem with Return route also.

```
9.|-- 209.85.244.25      0.0% 10 260.6 147.0 78.1 260.6 75.3
10.|-- ???              100.0 10 0.0 0.0 0.0 0.0 0.0
11.|-- 74.125.200.100    0.0% 10 84.8 112.4 75.6 234.4 63.9
```

In the above example, the router at hop 10 is either not responding to ICMP or there is a problem in the return route of the packet.

Telnet(TELecommunication NETwork)

- network protocol used on the Internet or local area network (LAN) connections
- bidirectional interactive communications facility.
- enables a user to connect to a remote host or device using a telnet client, usually over port 23.

telnet hostname

- **telnet hostname** would connect a user to a host named **hostname**.
- Telnet enables a user to manage an account or device remotely.
- For example, a user may telnet into a computer that hosts their website to manage his or her files remotely.
- telnet provides access to a command-line interface on a remote host via a virtual terminal connection
- which consists of an 8-bit byte oriented data connection over the Transmission Control Protocol (TCP).

The problem with telnet

Telnet sends everything in clear text format including username and password. You can use tcpdump or snoop to see all information.

telnet examples

telnet myhost.com

Attempts to open a connection to the remote host **myhost.com**. If a connection is established, the host will prompt for a login name and password.

telnet -l myusername myhost.com:5555

- Attempts to open a connection to the remote host **myhost.com** on port **5555**
- using the login name **myusername**.
- using the login name **myusername**.

telnet

- Opens a local **telnet>** prompt
- where you can enter any of the commands listed above.

telnet> open myhost.com

tcpdump

- **Tcpdump is a commandline network analyzer tool**
- **packet sniffer**
- which is used to capture or filter TCP/IP packets that received or transferred over a network on a specific interface.
- It is available under most of the **Linux/Unix** based operating systems.
- tcpdump also gives us a option to save captured packets in a file for future analysis.
- It saves the file in a **pcap** format
- that can be viewed by tcpdump command or a open source GUI based tool called Wireshark (Network Protocol Analyzer) that reads tcpdump **pcap** format files.

Installing in ubuntu

```
sudo apt-get install tcpdump
```

Redhat

```
Yum install tcpdump
```

Tcpdump depends on **libpcap library** for sniffing packets

Basic snipping

When you execute tcpdump command without any option, it will capture all the packets flowing through all the interfaces.

Basic sniffing

The first simple command to use is `tcpdump -n` (Don't convert addresses (i.e., host addresses, port numbers, etc.) to names.)

sudo tcpdump -n

tcpdump: verbose output suppressed, use -v or -vv for full protocol decode

listening on eth0, link-type EN10MB (Ethernet), capture size 65535 bytes

16:34:57.266865 IP 173.194.36.6 > 192.168.1.101: ICMP echo reply, id 19941, seq 1176, length 64

16:34:57.267226 IP 192.168.1.101.21271 > 218.248.255.163.53: 23380+ PTR? 6.36.194.173.in-addr.arpa. (43)

16:34:57.274549 IP 218.248.255.163.53 > 192.168.1.101.21271: 23380 1/4/2 PTR bom04s01-in-f6.1e100.net. (195)

16:34:57.297874 IP 192.168.1.101.56295 > 186.105.77.150.38213: UDP, length 105

- **Why sudo ?** Because tcpdump needs root privileges to be able to capture packets on network interfaces.
- The -n parameter is given to stop tcpdump from resolving ip addresses to hostnames
- 16:34:57.267226 IP 192.168.1.101.21271 > 218.248.255.163.53: 23380+ PTR? 6.36.194.173.in-addr.arpa. (43)
- "16:34:57.267226" is the timestamp with microsecond precision.
- **IP 192.168.1.101.21271** protocol of the packet called IP
- **218.248.255.163.53: 23380** source ip address joined with the source port
- Followed by destination ip
- **in-addr.arpa**-which is used to convert 32-bit numeric IP addresses back into domain names

verbose switch

```
sudo tcpdump -v -n
```

```
16:43:13.058660 IP (tos 0x20, ttl 54, id 50249, offset 0, flags [DF], proto TCP (6), length 40)
  64.41.140.209.5222 > 192.168.1.101.35783: Flags [.], cksum 0x6d32 (correct), ack 1617156745, win 9648, length 0
16:43:13.214621 IP (tos 0x0, ttl 64, id 0, offset 0, flags [DF], proto ICMP (1), length 84)
  192.168.1.101 > 173.194.36.6: ICMP echo request, id 19941, seq 1659, length 64
16:43:13.355334 IP (tos 0x20, ttl 54, id 48656, offset 0, flags [none], proto ICMP (1), length 84)
  173.194.36.6 > 192.168.1.101: ICMP echo reply, id 19941, seq 1659, length 64
```

- verbose switch lots of additional details about the packet
- include the **ttl**, **id**, **tcp flags**, **packet length** etc.

Getting the ethernet header

- In the above examples details of the ethernet header are not printed.
- **-e** option to print the ethernet header details

```
sudo tcpdump -vv -n -e
```

```
tcpdump: listening on eth0, link-type EN10MB (Ethernet), capture size 65535 bytes
17:57:27.218531 00:25:5e:1a:3d:f1 > 00:1c:c0:f8:79:ee, ethertype IPv4 (0x0800), length 98: (tos 0x20, ttl 54, id 53046, offset 0, flags [none], proto ICMP (1), length 84)
  173.194.36.6 > 192.168.1.101: ICMP echo reply, id 19941, seq 6015, length 64
```

Now the first thing after the timestamp is the source and destination mac address.

List of available interface

Sudo tcpdump -D

- 1.eth0
- 2.usbmon1 (USB bus number 1)
- 3.usbmon2 (USB bus number 2)
- 4.any (Pseudo-device that captures on all interfaces)
- 5.lo

Sniffing a particular interface

```
sudo tcpdump -i 1
```

```
sudo tcpdump -i eth0
```

Filtering packets using expressions

Selecting protocols

```
sudo tcpdump -n tcp
```

The above command will show only tcp packets. Similarly udp or icmp can be specified.

Particular host or port

```
sudo tcpdump -n 'src 192.168.1.101'
```

Search the network traffic using grep

```
$ sudo tcpdump -n -A | grep -e 'POST'
[sudo] password for enlightened:
tcpdump: verbose output suppressed, use -v or -vv for full protocol decode
listening on eth0, link-type EN10MB (Ethernet), capture size 65535 bytes
E...=.@.@.....e@.H..'.P(.o%~...P.9.PN..POST /blog/wp-admin/admin-ajax.php HTTP/1.1
E...c_@.@..=....e@.H..*.PfC<....wP.9.PN..POST /blog/wp-admin/admin-ajax.php HTTP/1.1
E.....@.@.....e@.H...."g;.(.-,WP.9.Nj..POST /login/?login_only=1 HTTP/1.1
```

-A Print each packet in ASCII.

sniffing password

```
tcpdump port http or port ftp or port smtp or port imap or port pop3 -l -A | egrep -i  
'pass=|pwd=|log=|login=|user=|username=|pw=|passw=|passwd=|password=|pass:|user:|username:|password:|login:|pass |user' --color=auto  
--line-buffered -B20
```

-l make stdout line buffered.useful if you want to see data while capturing it.

strace

- **strace** is a diagnostic, debugging and instructional [userspace](#) utility for [Linux](#).
- It is used to monitor interactions between [processes](#) and the [Linux kernel](#)
- which include [system calls](#), [signal](#) deliveries, and changes of process state
- The operation of strace is made possible by the kernel feature known as [ptrace](#).
- strace uses the ptrace system call to attach to the target process.
- When the target process makes a system call, it will be stopped by the kernel with a SIGTRAP signal and strace will be notified.
- CPU registers are the main interface for exchanging data between a process in the user space and the kernel
- so they are used to implement system calls
- strace will inspect the registers and stack of the target process
- and what the arguments were. strace then resumes the process.
- When it returns from the system call, it is stopped, and strace is notified again
- so it can inspect the return value. strace prints some information for the user each time this happens.
-

Usage

- The **most common usage is to start a program using strace, which prints a list of system calls made by the program.**
- This is useful if the program continually crashes, or does not behave as expected; for example using strace may reveal that the program is attempting to access a file which does not exist or cannot be read.
- An alternative application is to use the **-p** flag to attach to a running process.
- This is useful if a process has stopped responding, and might reveal, for example, that the process is blocking whilst attempting to make a network connection.
- As strace only details system calls, it cannot be used to detect as many problems as a code debugger

Trace the Execution of an Executable

strace command to trace the execution of any executable

```
strace ls

execve("/bin/ls", ["ls"], /* 21 vars */) = 0

brk(0)                      = 0x8c31000

access("/etc/ld.so.nohwcap", F_OK)      = -1 ENOENT (No such file or directory)

mmap2(NULL, 8192, PROT_READ, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) = 0xb78c7000

access("/etc/ld.so.preload", R_OK)      = -1 ENOENT (No such file or directory)

open("/etc/ld.so.cache", O_RDONLY)      = 3

fstat64(3, {st_mode=S_IFREG|0644, st_size=65354, ...}) = 0

...
```

```
getdents64(3, /* 0 entries */, 32768)  = 0
close(3)                                = 0
fstat64(1, {st_mode=S_IFCHR|0620, st_rdev=makedev(136, 0), ...}) = 0
mmap2(NULL, 4096, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) = 0xb7776000
write(1, "dead.letter Desktop Documents "..., 86dead.letter Desktop Documents Downloads Music Pictures Public Templates      Videos
) = 86
close(1)                                = 0
munmap(0xb7776000, 4096)                = 0
close(2)                                = 0
exit_group(0)                           = ?
```

The above part of the output shows the write system call where it outputs to STDOUT the current directory's listing.

Trace a Specific System Calls in an Executable Using Option -e

```
strace -e open ls
```

```
open("/etc/ld.so.cache", O_RDONLY)      = 3
open("/lib/libselinux.so.1", O_RDONLY)    = 3
open("/lib/librt.so.1", O_RDONLY)        = 3
open("/lib/libacl.so.1", O_RDONLY)        = 3
open("/lib/libc.so.6", O_RDONLY)          = 3
open("/lib/libdl.so.2", O_RDONLY)         = 3
open("/lib/libpthread.so.0", O_RDONLY)    = 3
open("/lib/libattr.so.1", O_RDONLY)        = 3
open("/proc/filesystems", O_RDONLY|O_LARGEFILE) = 3
open("/usr/lib/locale/locale-archive", O_RDONLY|O_LARGEFILE) = 3
open(".", O_RDONLY|O_NONBLOCK|O_LARGEFILE|O_DIRECTORY|O_CLOEXEC) = 3
```

```
Desktop  Documents  Downloads  examples.desktop  libflashplayer.so
```

- The above output displays **only the open system call** of the ls command.
- At the end of the strace output, it also displays the output of the ls command.

trace multiple system calls use the “-e trace=” option

```
$ strace -e trace=open,read ls /home
```

```
open("/etc/ld.so.cache", O_RDONLY)      = 3
open("/lib/libselinux.so.1", O_RDONLY)    = 3
read(3, "\177ELF\1\1\1\3\3\1\260G004"..., 512) = 512
open("/lib/librt.so.1", O_RDONLY)        = 3
read(3, "\177ELF\1\1\1\3\3\1\300\30004"..., 512) = 512
..
open("/lib/libattr.so.1", O_RDONLY)      = 3
read(3, "\177ELF\1\1\1\3\3\1\360\r004"..., 512) = 512
open("/proc/filesystems", O_RDONLY|O_LARGEFILE) = 3
read(3, "nodev\tsysfs\nnodev\trootfs\nnodev\tb"..., 1024) = 315
read(3, "", 1024)                      = 0
open("/usr/lib/locale/locale-archive", O_RDONLY|O_LARGEFILE) = 3
open("/home", O_RDONLY|O_NONBLOCK|O_LARGEFILE|O_DIRECTORY|O_CLOEXEC) = 3
```

Save the Trace Execution to a File Using Option -o

```
strace -o output.txt ls
```

```
Desktop  Documents  Downloads  examples.desktop  libflashplayer.so
```

```
Music  output.txt  Pictures  Public  Templates  Ubuntu_OS  Videos
```

```
cat output.txt
```

```
execve("/bin/ls", ["ls"], /* 37 vars */) = 0
brk(0) = 0x8637000
access("/etc/ld.so.nohwcap", F_OK) = -1 ENOENT (No such file or directory)
mmap2(NULL, 8192, PROT_READ, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) = 0xb7860000
access("/etc/ld.so.preload", R_OK) = -1 ENOENT (No such file or directory)
open("/etc/ld.so.cache", O_RDONLY) = 3
fstat64(3, {st_mode=S_IFREG|0644, st_size=67188, ...}) = 0
...
```

Execute Strace on a Running Linux Process Using Option -p

- execute strace on a program that is already running using the process id
- `ps -C firefox`

```
PID TTY          TIME CMD
1725 ?        00:40:50 firefox
```

```
sudo strace -p 1725 -o firefox_trace.txt
```

Print Timestamp for Each Trace Output Line Using Option -t

```
strace -t -e open ls /home
```

```
10:25:15 open("/etc/ld.so.cache", O_RDONLY|O_CLOEXEC) = 3
```

```
10:25:15 open("/lib/x86_64-linux-gnu/libselinux.so.1", O_RDONLY|O_CLOEXEC) = 3
```

```
10:25:15 open("/lib/x86_64-linux-gnu/libacl.so.1", O_RDONLY|O_CLOEXEC) = 3
```

```
10:25:15 open("/lib/x86_64-linux-gnu/libc.so.6", O_RDONLY|O_CLOEXEC) = 3
```

Print Relative Time for System Calls Using Option -r

Strace also has the option to print the execution time for each system calls

```
strace -r ls
```

```
strace -r ls

0.000000 execve("/bin/ls", ["ls"], /* 37 vars */) = 0
0.000846 brk(0)                      = 0x8418000
0.000143 access("/etc/ld.so.nohwcap", F_OK) = -1 ENOENT (No such file or directory)
0.000163 mmap2(NULL, 8192, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) = 0xb787b000
0.000119 access("/etc/ld.so.preload", R_OK) = -1 ENOENT (No such file or directory)
0.000123 open("/etc/ld.so.cache", O_RDONLY) = 3
```

Generate Statistics Report of System Calls Using Option -c

Using option -c, strace provides useful statistical report for the execution trace.

“calls” column in the following output indicated how many times that particular system call was executed.

% time	seconds	usecs/call	calls	errors	syscall
<hr/>					
23.62	0.000128	6	23		mmap
17.16	0.000093	7	14		mprotect
14.02	0.000076	8	9		open
7.56	0.000041	3	12		close
7.38	0.000040	5	8		read
7.38	0.000040	5	8	8	access
5.35	0.000029	10	3		munmap
4.80	0.000026	3	10		fstat
3.32	0.000018	9	2		getdents
3.32	0.000018	9	2	2	statfs
1.48	0.000008	8	1		write
1.11	0.000006	2	3		brk
1.11	0.000006	3	2		ioctl
0.92	0.000005	5	1		stat
0.74	0.000004	4	1		openat
0.37	0.000002	2	1		execve
0.37	0.000002	2	1		arch_prctl
<hr/>					
100.00	0.000542		101	10	total

Itrace

- ltrace lists all the library calls being called in an executable or a running process.
- Its name itself comes from **library-call tracing**.
- very useful for debugging user-space applications to **determine which library call is failing**.
- It is also capable of **receiving signals for segmentation faults**, etc.

Using ltrace

One does NOT need any root privileges to run ltrace.

Using ltrace to debug an application executable

ltrace ./executable <parameters>

Using ltrace to debug a running process

ltrace -p <PID>

Itrace time

```
__libc_start_main(0x8048a50, 1, 0xbfc806e4, 0x8049890, 0x8049880 <unfinished ...>
getenv("TIME")                                = NULL
 getopt_long(1, 0xbfc806e4, "+af:o:pqvV", 0x0804c0a0, NULL)           = -1
 __fprintf_chk(0x7b6980, 1, 0x8049948, 0xbfc81dce, 0Usage: time [-apvV] [-f format] [-o file] [--append] [--verbose]
 [-portability] [--format=format] [--output=file] [--version]
 [--quiet] [--help] command [arg...]
)                                              = 177
exit(1 <unfinished ...>
+++ exited (status 1) +++
```

Debugging with ltrace

- ltrace utility is most useful for debugging user-space application programs
- as the bugs could be surfaced out seeing the dynamic records of the library calls.

- Lets take a sample program.

- This program has a bug,

- 1. **#include <stdio.h>**

- 2. **#include <unistd.h>**

- 3.

- 4. **int main()**

- 5. **{**

- 6. **FILE *fp = fopen("rfile.txt", "w+");**

- 7. **fprintf(fp+1, "Invalid Write\n");**

- 8. **fclose(fp);**

- 9. **return 0;**

- 10. **}**

-

Lets compile and run it.

```
gcc file.c -Wall -o file
```

```
./file
```

```
Segmentation fault (core dumped)
```

Lets ltrace the file and see whats happening.

```
ltrace ./file
```

```
__libc_start_main(0x8048454, 1, 0xbfc19db4, 0x80484c0, 0x8048530 <unfinished ...>
fopen("rfile.txt", "w+") = 0x9160008
fwrite("Invalid Write\n", 1, 14, 0x916009c <unfinished ...>
--- SIGSEGV (Segmentation fault) ---
+++ killed by SIGSEGV +++
```

- The first line states the libc function call 'main()', with its parameters and it terminated unfinished.
- Next line has listed the fopen() call to open 'rfile.txt' in write plus mode and returning a valid file pointer address.
- Moving on to the next library call, we have fwrite which is unfinished.
- Here is the bug because of which program crashed.
- notice the file pointer containing the address in fwrite() call doesnt match the one in fopen() call.
- The program is trying to write onto a file using file pointer which is invalid following a SIGSEGV signal.

- the bug is in line 7 of our program where the file pointer being passed to the fprintf() call is tampered. So, we just found the bug.

Corrected program

```
#include <stdio.h>
2. #include <unistd.h>
3.
4. int main()
5. {
6.     FILE *fp = fopen("rfile.txt", "w+");
7.     fprintf(fp, "Invalid Write\n");
8.     fclose(fp);
9.     return 0;
10. }
```

trace output would had been much efficient to debug than re-checking each line of source or adding various printf's.

/proc

- It doesn't really exist
- Its zero-length files are neither binary nor text
- it is a **virtual file system**.
- the operating system creates them on the fly if you try to read them.
- This special directory holds all the details about your Linux system
- including its kernel, processes, and configuration parameters.
- Most virtual files always have a current timestamp, which indicates that they are constantly being kept up to date.
- The /proc directory itself is created every time you boot your box.
- Need to be root to examine the whole directory
- /proc directory, you'll see two types of content — numbered directories, and system information files.
- Several Linux commands access the information from /proc, and displays in a certain format.

ls /proc

```
1 2432 3340 3715 3762 5441 815 devices modules
129 2474 3358 3716 3764 5445 acpi diskstats mounts
1290 248 3413 3717 3812 5459 asound dma mtrr
133 2486 3435 3718 3813 5479 bus execdomains partitions
1420 2489 3439 3728 3814 557 dri fb self
165 276 3450 3731 39 5842 driver filesystems slabinfo
166 280 36 3733 3973 5854 fs interrupts splash
2 2812 3602 3734 4 6 ide iomem stat
2267 3 3603 3735 40 6381 irq ioports swaps
2268 326 3614 3737 4083 6558 net kallsyms sysrq-trigger
2282 327 3696 3739 4868 6561 scsi kcore timer_list
2285 3284 3697 3742 4873 6961 sys keys timer_stats
2295 329 3700 3744 4878 7206 sysvipc key-users uptime
2335 3295 3701 3745 5 7207 tty kmsg version
2400 330 3706 3747 5109 7222 buddyinfo loadavg vmcore
2401 3318 3709 3749 5112 7225 cmdline locks vmstat
2427 3329 3710 3751 541 7244 config.gz meminfo zoneinfo
2428 3336 3714 3753 5440 752 cpuinfo misc
```

- lot of directories with just numbers.
- These numbers represents the process ids
- the files inside this numbered directory corresponds to the process with that particular PID

Some Important files under numbered directory

- cmdline – command line of the command.
- environ – environment variables.
- fd – Contains the file descriptors which is linked to the appropriate files
- limits – Contains the information about the specific limits to the process.
- mounts – mount related information

Important links under each numbered directory

- cwd – Link to current working directory of the process.
- exe – Link to executable of the process.
- root – Link to the root directory of the process.
-

/proc Files about the system information

- /proc/cpuinfo – information about CPU,
- /proc/meminfo – information about memory,
- /proc/loadvg – load average,
- /proc/partitions – partition related information,
- /proc/version – linux version

Some Linux commands read the information from this /proc files and displays it. For example, **free command**, reads the memory information from **/proc/meminfo file**, formats it, and displays it.

Example cd /proc/12

cat /proc/12/status

Name: xenwatch
State: S (sleeping)
Tgid: 12
Pid: 12
PPid: 2
TracerPid: 0
Uid: 0 0 0 0
Gid: 0 0 0 0
FDSize:64
Groups:
Threads: 1
SigQ: 1/4592
SigPnd: 0000000000000000
ShdPnd: 0000000000000000
SigBlk: 0000000000000000
SigIgn: ffffffff
SigCgt: 0000000000000000
CapInh: 0000000000000000
CapPrm: ffffffff
CapEff: ffffffff
CapBnd: ffffffff
Cpus_allowed: 1
Cpus_allowed_list: 0
Mems_allowed: 00000000,00000001
Mems_allowed_list: 0
voluntary_ctxt_switches: 84

- this process belongs to **xenwatch**
- Its current state is **sleeping**
- its process **ID** is **12**
- We also can see who is running this, as **UID** and **GID** are **0**, indicating that this process belongs to the **root** user.
-

free

- 'free' shows you information about the machine's memory.
- This includes physical memory (RAM), swap as well as the shared memory and buffers used by the kernel.
- All measurements are in Kb.

free

	total	used	free	shared	buffers	cached
Mem:	7993580	4765408	3228172	640200	189588	2242860
-/+ buffers/cache:	2332960	5660620				
Swap:	0	0	0			

total: total ram.

used: This is a mix of application used memory and other 'temporarily' (buffer + cache) used memory that is actually available if needed.

The 'temporarily' used memory is borrowed if available by the linux system to help speed up system performance,
Much of this type of memory is shown under the 'cached' column.

This memory is given up by the linux system if an application need memory.

free: free and untouched memory.

shared: Memory specifically allocated for use by multiple processes

buffers: Temporary memory that is set aside to help some processes

cache: Memory that is available and 'borrowed' by the operating system to help speed up many linux OS operations. This memory is given up by the system if an application need it.

-/+ buffers/cache - If you're wanting to get a good idea on how **much free memory is available**, the **free** section in the **buffers/cache** is what you should be reading.

-/+ buffers/cache is typically more helpful than the first **Mem** line.

The **intersection of free** and **-/+ buffers/cache** is essentially what you have for '**available**' **memory**.

Swap - memory management involving swapping regions of memory to and from storage.

Display Memory in Bytes

free -b

	total	used	free	shared	buffers	cached
Mem:	1046147072	934420480	111726592	0	123256832	671281152
-/+ buffers/cache:	139882496	906264576				
Swap:	4294959104	0	4294959104			

Display Memory in Kilo Bytes

free -k

	total	used	free	shared	buffers	cached
Mem:	1021628	912520	109108	0	120368	655548
-/+ buffers/cache:	136604	885024				
Swap:	4194296	0	4194296			

Display Memory in Megabytes

free -m

	total	used	free	shared	buffers	cached
Mem:	997	891	106	0	117	640
-/+ buffers/cache:	133	864				
Swap:	4095	0	4095			

Display Total Line

Free command with **-t** option, will list the total line at the end.

```
free -t
```

	total	used	free	shared	buffers	cached
Mem:	7993580	5188684	2804896	624840	200980	2372464
-/+ buffers/cache:	2615240	5378340				
Swap:	0	0	0			
Total:	7993580	5188684	2804896			

Disable Display of Buffer Adjusted Line

```
free -o
```

	total	used	free	shared	buffers	cached
Mem:	1021628	912520	109108	0	120368	655548
Swap:	4194296	0	4194296			

Display Memory Status for Regular Intervals

the below command will update free command every **3 seconds**.

free -s 3

	total	used	free	shared	buffers	cached
Mem:	7993580	5169700	2823880	639732	201472	2387352
-/+ buffers/cache:	2580876	5412704				
Swap:	0	0	0			

	total	used	free	shared	buffers	cached
Mem:	7993580	5165896	2827684	639404	201484	2387024
-/+ buffers/cache:	2577388	5416192				
Swap:	0	0	0			

	total	used	free	shared	buffers	cached
Mem:	7993580	5165468	2828112	639404	201484	2387024
-/+ buffers/cache:	2576960	5416620				
Swap:	0	0	0			

Show Low and High Memory Statistics

The **-l** switch displays detailed high and low memory size statistics.

free -l

	total	used	free	shared	buffers	cached
Mem:	1021628	912368	109260	0	120368	655548
Low:	890036	789064	100972			
High:	131592	123304	8288			
-/+ buffers/cache:	136452	885176				
Swap:	4194296	0	4194296			

top

- The top command in Linux **displays the running processes on the system.**
- It is used extensively for **monitoring the load on a server.**
- The top command is an interactive command
- **top** provides an ongoing look at **processor activity in real time.**
- displays the top 30 processes on the system and periodically updates this information.

```
$ top
```

raghu@raghu-Inspiron-1440: ~

File Edit View Search Terminal Help

top - 00:11:08 up 1:18, 2 users, load average: 0.22, 0.50, 0.55
Tasks: 149 total, 2 running, 147 sleeping, 0 stopped, 0 zombie
%Cpu(s): 2.7 us, 1.8 sy, 0.0 ni, 94.0 id, 1.5 wa, 0.0 hi, 0.0 si, 0.0 st
KiB Mem: 3065928 total, 1186836 used, 1879092 free, 111452 buffers
KiB Swap: 4095996 total, 0 used, 4095996 free, 675916 cached

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
1713	root	20	0	82832	24m	9.9m	S	4.0	0.8	1:12.95	Xorg
31435	raghu	20	0	387m	51m	18m	S	2.3	1.7	0:05.59	shutter
2260	raghu	20	0	174m	13m	9656	S	0.7	0.4	0:04.72	marco
31433	raghu	20	0	5208	1384	1032	R	0.7	0.0	0:00.44	top
10	root	20	0	0	0	0	S	0.3	0.0	0:10.22	ksoftirqd/1
1511	mysql	20	0	312m	38m	5928	S	0.3	1.3	0:03.43	mysqld
1892	root	20	0	120m	118m	608	D	0.3	4.0	0:02.69	HWActivator
2285	raghu	20	0	171m	11m	9232	S	0.3	0.4	0:00.46	mate-power-mana
2327	root	20	0	29012	4272	3200	S	0.3	0.1	0:02.45	upowerd
30049	root	20	0	0	0	0	S	0.3	0.0	0:00.26	kworker/1:3
31208	raghu	20	0	192m	13m	10m	S	0.3	0.5	0:00.56	mate-terminal
1	root	20	0	3728	2172	1348	S	0.0	0.1	0:01.10	init
2	root	20	0	0	0	0	S	0.0	0.0	0:00.00	kthreadd
3	root	20	0	0	0	0	S	0.0	0.0	0:09.98	ksoftirqd/0
6	root	rt	0	0	0	0	S	0.0	0.0	0:00.27	migration/0
7	root	rt	0	0	0	0	S	0.0	0.0	0:00.06	watchdog/0
8	root	rt	0	0	0	0	S	0.0	0.0	0:00.10	migration/1

Uptime and Load Averages:

```
top - 00:15:48 up 1:23, 2 users, load average: 0.26, 0.33, 0.47
```

- output similar to uptime command.
- The fields display:
 - current time
 - the time your system is been up
 - number of users logged in
 - load average of 5, 10 and 15 minutes respectively.

Tasks:

```
Tasks: 147 total, 2 running, 145 sleeping, 0 stopped, 0 zombie
```

- The second line shows summary of tasks or processes. The processes can be in different states.
- It shows total number of the processes
- Out of these, the processes can be running, sleeping, stopped or in zombie

CPU States:

```
%Cpu(s): 1.8 us, 1.2 sy, 0.0 ni, 97.0 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
```

% of CPU time in different modes is shown. The meaning of different CPU times are:

- * **us**, user: CPU time in running (un-niced) user processes
- * **sy**, system: CPU time in running kernel processes
- * **ni**, niced: CPU time in running niced user processes
- * **wa**, IO wait: CPU time waiting for IO completion
- * **hi**: CPU time serving hardware interrupts
- * **si**: CPU time serving software interrupts
- * **st**: CPU time stolen for this vm by the hypervisor.

Memory Usage:

```
KiB Mem: 3065928 total, 1444436 used, 1621492 free, 189296 buffers
KiB Swap: 4095996 total, 0 used, 4095996 free, 766820 cached
```

- two lines show memory usage
- somewhat like 'free' command.
- 1st of these lines is for **physical memory** and the second for **virtual memory** (swap space).

- The physical memory is displayed as: total available memory, used memory, free memory, and memory used for buffers
- Similarly, swap reflects: total, used, free and cached swap space.

Fields/Columns:

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
31435	raghu	20	0	401m	65m	19m	S	7.3	2.2	1:08.82	shutter
1713	root	20	0	84176	25m	10m	S	4.0	0.9	2:00.70	Xorg
2490	raghu	20	0	182m	13m	10m	S	2.0	0.4	0:05.58	wnck-applet
2260	raghu	20	0	175m	14m	10m	S	0.7	0.5	0:10.90	marco
2228	raghu	20	0	101m	11m	9072	S	0.3	0.4	0:01.52	mate-settings-d
2262	raghu	20	0	197m	18m	11m	S	0.3	0.6	0:04.45	mate-panel
2284	raghu	20	0	222m	32m	20m	S	0.3	1.1	0:16.13	caja
27689	raghu	20	0	202m	29m	16m	S	0.3	1.0	0:09.58	pluma
31208	raghu	20	0	192m	13m	10m	S	0.3	0.5	0:01.97	mate-terminal
31433	raghu	20	0	5208	1444	1080	R	0.3	0.0	0:04.45	top
1	root	20	0	3728	2172	1348	S	0.0	0.1	0:01.14	init
2	root	20	0	0	0	0	S	0.0	0.0	0:00.00	kthreadd
3	root	20	0	0	0	0	S	0.0	0.0	0:12.94	ksoftirqd/0
6	root	rt	0	0	0	0	S	0.0	0.0	0:00.27	migration/0
7	root	rt	0	0	0	0	S	0.0	0.0	0:00.07	watchdog/0
8	root	rt	0	0	0	0	S	0.0	0.0	0:00.10	migration/1
10	root	20	0	0	0	0	S	0.0	0.0	0:12.64	ksoftirqd/1

PID

The Process ID, to uniquely identify a processes.

USER

The effective user name of the owner of the processes.

PR

The **scheduling priority of the process**. Some values in this field are 'rt'. It means that the **process is running under real-time**.

NI

The **nice value of the process**. Lower values mean higher priority.

VIRT

The amount of **virtual memory used by the process**.

RES

The **resident memory size**. Resident memory is the amount of non-swapped **physical memory a task is using**.

SHR

SHR is the **shared memory** used by the process.

S

This is the process status. It can have one of the following values:

D - uninterruptible sleep

R - running

S - sleeping

T - traced or stopped

Z - zombie

%CPU

It is the **percentage of CPU time the task has used** since last update.

%MEM

Percentage of **available physical memory used by the process**.

TIME+

The **total CPU time the task has used since it started**, with precision upto hundredth of a second.

COMMAND

The command which was used to start the process.

Display Specific User Process

Use top command with 'u' option will display specific **User** process details.

```
top -u nandaraj
```

Tasks: 230 total, 4 running, 226 sleeping, 0 stopped, 0 zombie
 %Cpu(s): 0.9 us, 0.2 sy, 0.0 ni, 98.8 id, 0.1 wa, 0.0 hi, 0.0 si, 0.0 st
 KiB Mem: 7993580 total, 5502568 used, 2491012 free, 212404 buffers
 KiB Swap: 0 total, 0 used, 0 free. 2542400 cached Mem

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
3105	nandaraj	20	0	1941256	977448	60268	R	6.6	12.2	47:36.16	firefox
2409	nandaraj	20	0	1559312	206280	36220	S	1.7	2.6	4:50.97	compiz
4638	nandaraj	20	0	630756	19732	12532	S	1.7	0.2	0:04.70	gnome-terminal
2919	nandaraj	20	0	1143964	249132	40916	S	0.3	3.1	1:11.51	chrome
2116	nandaraj	20	0	600304	6724	3520	S	0.0	0.1	0:00.24	gnome-keyring
2131	nandaraj	20	0	40188	2628	1604	S	0.0	0.0	0:00.15	init
2218	nandaraj	20	0	40336	2524	928	S	0.0	0.0	0:01.71	dbus-daemon
2229	nandaraj	20	0	22296	1156	992	S	0.0	0.0	0:00.00	upstart-event
2234	nandaraj	20	0	78320	4356	3768	S	0.0	0.1	0:00.12	window-starter
2253	nandaraj	20	0	365892	9496	2892	S	0.0	0.1	0:16.43	ibus-daemon
2278	nandaraj	20	0	30780	712	408	S	0.0	0.0	0:00.00	upstart-finder
2280	nandaraj	20	0	22388	724	424	S	0.0	0.0	0:00.31	upstart-dbus
2282	nandaraj	20	0	22304	640	412	S	0.0	0.0	0:00.02	upstart-dbus
2296	nandaraj	20	0	813368	17856	11508	S	0.0	0.2	0:00.50	unity-settings
2300	nandaraj	20	0	647900	28028	15892	S	0.0	0.4	0:00.99	hud-service
2303	nandaraj	20	0	337576	3428	2880	S	0.0	0.0	0:00.00	at-spi-bus
2304	nandaraj	20	0	583828	11296	7796	S	0.0	0.1	0:00.19	gnome-session

nandaraj@DESKTOP-3348S: ~

Highlight Running Process in Top

Press 'z' option in running top command will display running process in color which may help you to identified running process easily.

top - 12:12:49 up 4:04, 2 users, load average: 0.75, 0.36, 0.35											
Tasks: 234 total, 1 running, 233 sleeping, 0 stopped, 0 zombie											
%Cpu(s): 3.4 us, 0.9 sy, 0.0 ni, 95.2 id, 0.4 wa, 0.0 hi, 0.0 si, 0.0 st											
KiB Mem: 7993580 total, 5511612 used, 2481968 free, 213008 buffers											
KiB Swap: 0 total, 0 used, 0 free. 2539760 cached Mem											
PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
3105	nandaraj	20	0	1938696	970408	60268	S	6.0	12.1	48:27.94	firefox
2409	nandaraj	20	0	1560096	206848	36740	S	5.0	2.6	4:55.20	compiz
1689	root	20	0	700908	123056	102604	S	3.7	1.5	8:24.18	Xorg
6066	nandaraj	20	0	609616	16368	11640	S	2.7	0.2	0:00.24	gnome-scre
7	root	20	0	0	0	0	S	0.3	0.0	0:03.12	rcu_sched
11	root	20	0	0	0	0	S	0.3	0.0	0:00.66	rcuos/3
214	root	20	0	0	0	0	S	0.3	0.0	0:01.87	kworker/3:1
2663	nandaraj	20	0	1025412	123480	61532	S	0.3	1.5	0:51.08	chrome
6052	nandaraj	20	0	29144	1740	1176	R	0.3	0.0	0:00.09	top
1	root	20	0	33908	3308	1500	S	0.0	0.0	0:00.93	init
2	root	20	0	0	0	0	S	0.0	0.0	0:00.00	kthreadd
3	root	20	0	0	0	0	S	0.0	0.0	0:00.04	ksoftirqd/0
5	root	0	-20	0	0	0	S	0.0	0.0	0:00.00	kworker/0:+
8	root	20	0	0	0	0	S	0.0	0.0	0:02.52	rcuos/0
9	root	20	0	0	0	0	S	0.0	0.0	0:00.78	rcuos/1
10	root	20	0	0	0	0	S	0.0	0.0	0:00.63	rcuos/2
12	root	20	0	0	0	0	S	0.0	0.0	0:01.14	rcuos/4

Shows Absolute Path of Processes

Press 'c' option in running top command, it will display absolute path of running process.

Tasks: 232 total, 4 running, 228 sleeping, 0 stopped, 0 zombie										
%Cpu(s): 1.2 us, 0.3 sy, 0.0 ni, 98.5 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st										
KiB Mem: 7993580 total, 5511584 used, 2481996 free, 213368 buffers										
KiB Swap: 0 total, 0 used, 0 free. 2541124 cached Mem										
PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+ COMMAND
3105	nandaraj	20	0	1936648	987304	60268	R	4.7	12.4	48:50.45 firefox
2409	nandaraj	20	0	1559320	206344	36284	S	1.8	2.6	4:57.32 compiz
4638	nandaraj	20	0	632036	19960	12556	S	1.8	0.2	0:05.60 gnome-terminal
1689	root	20	0	697368	127236	105752	S	0.6	1.6	8:26.48 Xorg
2253	nandaraj	20	0	366152	9756	2892	R	0.6	0.1	0:17.28 ibus-daemon

```
Tasks: 231 total, 1 running, 230 sleeping, 0 stopped, 0 zombie
%Cpu(s): 2.6 us, 0.1 sy, 0.0 ni, 97.3 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
KiB Mem: 7993580 total, 5516932 used, 2476648 free, 213540 buffers
KiB Swap: 0 total, 0 used, 0 free. 2521936 cached Mem
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+ COMMAND
3105	nandaraj	20	0	1943880	996556	60268	S	20.8	12.5	49:01.74 /usr/lib/firefox/firefox
4638	nandaraj	20	0	631092	19824	12556	S	1.3	0.2	0:05.85 gnome-terminal
75	root	20	0	0	0	0	S	0.7	0.0	0:01.95 [kworker/1:1]
2253	nandaraj	20	0	366152	9756	2892	S	0.7	0.1	0:17.55 /usr/bin/ibus-daemon --daemonize --xim
2387	nandaraj	20	0	205148	3284	2708	S	0.7	0.0	0:03.69 /usr/lib/ibus/ibus-engine-simple

Set 'Screen Refresh Interval' in Top

By default screen refresh interval is **3.0** seconds, same can be changed by pressing '**d**' option in running top command

```
top - 12:20:33 up 4:11, 2 users, load average: 0.56, 0.30, 0.32
Tasks: 231 total, 1 running, 230 sleeping, 0 stopped, 0 zombie
%Cpu(s): 2.8 us, 0.4 sy, 0.0 ni, 96.2 id, 0.5 wa, 0.0 hi, 0.0 si, 0.0 st
KiB Mem: 7993580 total, 5504212 used, 2489368 free, 214156 buffers
KiB Swap: 0 total, 0 used, 0 free. 2555508 cached Mem
Change delay from 3.0 to 1.0
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
2253	nandaraj	20	0	366152	9756	2892	S	6.4	0.1	0:17.91	ibus-daemon
1	root	20	0	33908	3308	1500	S	0.0	0.0	0:00.93	init

Kill running process with argument 'k'

You can kill a process after finding **PID** of process by pressing '**k**' option in running top command without exiting from top window as shown below.

```
top - 12:22:24 up 4:13, 2 users, load average: 0.65, 0.42, 0.36
Tasks: 231 total, 2 running, 229 sleeping, 0 stopped, 0 zombie
%Cpu(s): 4.2 us, 1.0 sy, 0.0 ni, 94.7 id, 0.1 wa, 0.0 hi, 0.0 si, 0.0 st
KiB Mem: 7993580 total, 5514344 used, 2479236 free, 214440 buffers
KiB Swap: 0 total, 0 used, 0 free. 2556248 cached Mem
PID to signal/kill [default pid = 3105] 2409
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
3105	nandaraj	20	0	1947336	965764	60268	R	28.6	12.1	50:24.36	firefox
1689	root	20	0	714236	127284	105800	S	2.7	1.6	8:37.87	Xorg
2409	nandaraj	20	0	1559572	206344	36284	S	2.0	2.6	5:05.10	compiz

Sort by CPU Utilisation

Press (**Shift+P**) to sort processes as per CPU utilization.

```
Tasks: 232 total, 1 running, 231 sleeping, 0 stopped, 0 zombie
%Cpu(s): 0.6 us, 0.1 sy, 0.0 ni, 99.2 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
KiB Mem: 7993580 total, 5530884 used, 2462696 free, 214752 buffers
KiB Swap: 0 total, 0 used, 0 free. 2549228 cached Mem
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
3105	nandaraj	20	0	1952200	989724	60268	S	5.6	12.4	50:54.84	firefox
1689	root	20	0	738720	142444	120960	S	0.3	1.8	8:41.49	Xorg
2409	nandaraj	20	0	1559708	206344	36284	S	0.3	2.6	5:07.06	compiz
4638	nandaraj	20	0	629652	19456	12556	S	0.3	0.2	0:06.59	gnome-term+
6170	nandaraj	20	0	924824	21020	13388	S	0.3	0.3	0:00.75	gnome-scre+
1	root	20	0	33908	3308	1500	S	0.0	0.0	0:00.93	init
2	root	20	0	0	0	0	S	0.0	0.0	0:00.00	kthreadd
3	root	20	0	0	0	0	S	0.0	0.0	0:00.04	ksoftirqd/0
5	root	0	-20	0	0	0	S	0.0	0.0	0:00.00	kworker/0:+
7	root	20	0	0	0	0	S	0.0	0.0	0:03.30	rcu_sched
8	root	20	0	0	0	0	S	0.0	0.0	0:02.70	rcuos/0
9	root	20	0	0	0	0	S	0.0	0.0	0:00.00	l1

Exit Top Command After Specific repetition

Top output keep refreshing until you press 'q'. With below command top command will automatically exit after 10 number of repetition.

```
# top -n 10
```

Getting Top Command Help

Press 'h' option to obtain the top command help.

```
Help for Interactive Commands - procps-ng version 3.3.9
Window 1:Def: Cumulative mode Off.  System: Delay 3.0 secs; Secure mode off.

Z,B,E,e  Global: 'Z' colors; 'B' bold; 'E'/'e' summary/task memory scale
l,t,m  Toggle Summary: 'l' load avg; 't' task/cpu stats; 'm' memory info
0,1,2,3,I  Toggle: '0' zeros; '1/2/3' cpus or numa node views; 'I' Irix mode
f,F,X  Fields: 'f'/'F' add/remove/order/sort; 'X' increase fixed-width

L,&,<,>  Locate: 'L'/'&' find/again; Move sort column: '<'/'>' left/right
R,H,V,J  Toggle: 'R' Sort; 'H' Threads; 'V' Forest view; 'J' Num justify
c,i,S,j  Toggle: 'c' Cmd name/line; 'i' Idle; 'S' Time; 'j' Str justify
x,y  Toggle highlights: 'x' sort field; 'y' running tasks
z,b  Toggle: 'z' color/mono; 'b' bold/reverse (only if 'x' or 'y')
u,U,o,O  Filter by: 'u'/'U' effective/any user; 'o'/'O' other criteria
n,#,^0  Set: 'n'/'#' max tasks displayed; Show: Ctrl+'0' other filter(s)
C,...  Toggle scroll coordinates msg for: up,down,left,right,home,end

k,r  Manipulate tasks: 'k' kill; 'r' renice
d or s  Set update interval
w,Y  Write configuration file 'W'; Inspect other output 'Y'
q  Quit
( commands shown with '.' require a visible task display window )
Press 'h' or '?' for help with Windows,
Type 'q' or <Esc> to continue █
```

ps

ps

- **ps** program (short for "process **s**tatus") **displays the currently-running processes**

- **ps**

- ```
root@HIFX-33:~# ps
 PID TTY TIME CMD
 6232 pts/14 00:00:00 sudo
 6233 pts/14 00:00:00 bash
 7089 pts/14 00:00:00 ps
```

**PID** - The process ID of the process.

**TTY** The **controlling terminal** for the process.

**TIME** - The **CPU time the process has used since the process started**. The format for this time is [ dd-]hh:mm:ss.

**CMD** - The name of the command that is running

## Display User-Oriented Output (u)

ps u

| USER     | PID  | %CPU | %MEM | VSZ   | RSS  | TTY    | STAT | START | TIME | COMMAND |
|----------|------|------|------|-------|------|--------|------|-------|------|---------|
| hendaraj | 4648 | 0.0  | 0.0  | 28468 | 5596 | pts/14 | Ss   | 08:50 | 0:00 | bash    |
| hendaraj | 6256 | 0.0  | 0.0  | 28352 | 5324 | pts/0  | Ss   | 12:31 | 0:00 | bash    |
| hendaraj | 7137 | 0.0  | 0.0  | 22640 | 1312 | pts/0  | R+   | 14:18 | 0:00 | ps u    |

**USER**-The user that the process is running under.

**%CPU**-The **percentage of time the process has used the CPU** since the process started.

**%MEM**-The percentage of **real memory used** by this process.

**VSZ**-Indicates, as a decimal integer, the **size in kilobytes of the process in virtual memory**.

**RSS**-The **real-memory size of the process** (in 1 KB units).

**STAT**-The status of the process.

**STARTED**-When the process was started.

# Display All Processes (-e)

Display all processes in the system irrespective of user.

mix of “system” processes as well as “user” processes.

```
ps -e
```

**Listing 6.3**

| PID   | TTY     | TIME     | CMD                                                     |
|-------|---------|----------|---------------------------------------------------------|
| 1     | ??      | 28:22.66 | /sbin/launchd                                           |
| 17    | ??      | 1:19.09  | /usr/libexec/UserEventAgent (System)                    |
| 18    | ??      | 0:17.57  | /usr/libexec/kextd                                      |
| 19    | ??      | 0:36.91  | /usr/libexec/taskgated -s                               |
| 20    | ??      | 0:25.96  | /usr/sbin/notifyd                                       |
| 21    | ??      | 0:19.43  | /usr/sbin/securityd -i                                  |
| 22    | ??      | 0:17.15  | /usr/libexec/diskarbitrationd                           |
| 23    | ??      | 0:21.38  | /System/Library/CoreServices/powerd.bundle/powerd       |
| 24    | ??      | 5:13.79  | /usr/libexec/configd                                    |
| 25    | ??      | 1:01.05  | /usr/sbin/syslogd                                       |
| ...   |         |          |                                                         |
| 24703 | ttys004 | 0:00.15  | -bash                                                   |
| 24743 | ttys005 | 0:00.02  | login -pfl markbates /bin/bash -c exec -la bash /bin... |
| 24755 | ttys005 | 0:00.21  | -bash                                                   |
| 57892 | ttys005 | 0:00.04  | ruby /Users/markbates/scripts/ss                        |
| 57896 | ttys005 | 0:00.42  | forego start -p 3000                                    |
| 24810 | ttys006 | 0:00.02  | login -pfl markbates /bin/bash -c exec -la bash /bin... |
| 24851 | ttys006 | 0:00.70  | -bash                                                   |
| 25864 | ttys007 | 0:00.01  | login -pfl markbates /bin/bash -c exec -la bash /bin... |
| 25865 | ttys007 | 0:00.24  | -bash                                                   |
| 25866 | ttys008 | 0:00.02  | login -pfl markbates /bin/bash -c exec -la bash /bin... |
| 25876 | ttys008 | 0:00.10  | -bash                                                   |

## Display Processes by User (-U)

- filter processes by the user that is running them.
- ps -U root

| PID | TTY | TIME     | CMD          |
|-----|-----|----------|--------------|
| 1   | ?   | 00:00:00 | init         |
| 2   | ?   | 00:00:00 | kthreadd     |
| 3   | ?   | 00:00:00 | ksoftirqd/0  |
| 5   | ?   | 00:00:00 | kworker/0:0H |
| 7   | ?   | 00:00:05 | rcu_sched    |
| 8   | ?   | 00:00:04 | rcuos/0      |
| 9   | ?   | 00:00:01 | rcuos/1      |
| 10  | ?   | 00:00:01 | rcuos/2      |
| 11  | ?   | 00:00:01 | rcuos/3      |
| 12  | ?   | 00:00:01 | rcuos/4      |
| 13  | ?   | 00:00:01 | rcuos/5      |
| 14  | ?   | 00:00:00 | rcuos/6      |
| 15  | ?   | 00:00:00 | rcuos/7      |
| 16  | ?   | 00:00:00 | rcu_bh       |
| 17  | ?   | 00:00:00 | rcuob/0      |
| 18  | ?   | 00:00:00 | rcuob/1      |

## ps -ef

To find out information about *all* the processes in the system, the command is,

| UID  | PID | PPID | C | STIME | TTY | TIME     | CMD            |
|------|-----|------|---|-------|-----|----------|----------------|
| root | 1   | 0    | 0 | 08:08 | ?   | 00:00:00 | /sbin/init     |
| root | 2   | 0    | 0 | 08:08 | ?   | 00:00:00 | [kthreadd]     |
| root | 3   | 2    | 0 | 08:08 | ?   | 00:00:00 | [ksoftirqd/0]  |
| root | 5   | 2    | 0 | 08:08 | ?   | 00:00:00 | [kworker/0:0H] |
| root | 7   | 2    | 0 | 08:08 | ?   | 00:00:05 | [rcu_sched]    |
| root | 8   | 2    | 0 | 08:08 | ?   | 00:00:04 | [rcuos/0]      |
| root | 9   | 2    | 0 | 08:08 | ?   | 00:00:01 | [rcuos/1]      |
| root | 10  | 2    | 0 | 08:08 | ?   | 00:00:01 | [rcuos/2]      |
| root | 11  | 2    | 0 | 08:08 | ?   | 00:00:01 | [rcuos/3]      |
| root | 12  | 2    | 0 | 08:08 | ?   | 00:00:01 | [rcuos/4]      |
| root | 13  | 2    | 0 | 08:08 | ?   | 00:00:01 | [rcuos/5]      |
| root | 14  | 2    | 0 | 08:08 | ?   | 00:00:00 | [rcuos/6]      |
| root | 15  | 2    | 0 | 08:08 | ?   | 00:00:00 | [rcuos/7]      |
| root | 16  | 2    | 0 | 08:08 | ?   | 00:00:00 | [rcu_bh]       |

- The **-e** option selects all processes in the system.
- Without this, ps gives the processes belonging to the same effective user id
- **The -f option gives the full listing.**
- **UID** - gives the process owner's user-id
- **PID** - process-id
- **PPID** - parent process id

- **C** - represents CPU utilization.
- A high value for C indicates a CPU intensive process.
- **STIME**- process start time
- TTY - terminal from which the process has been started
- TIME -cumulative **CPU time** used by the process in [DD-]hh:mm:ss format
- CMD - command, using which, the process was created

- **ps aux**

| USER | PID | %CPU | %MEM | VSZ   | RSS  | TTY | STAT | START | TIME | COMMAND        |
|------|-----|------|------|-------|------|-----|------|-------|------|----------------|
| root | 1   | 0.0  | 0.0  | 33908 | 3308 | ?   | Ss   | 08:08 | 0:00 | /sbin/init     |
| root | 2   | 0.0  | 0.0  | 0     | 0    | ?   | S    | 08:08 | 0:00 | [kthreadd]     |
| root | 3   | 0.0  | 0.0  | 0     | 0    | ?   | S    | 08:08 | 0:00 | [ksoftirqd/0]  |
| root | 5   | 0.0  | 0.0  | 0     | 0    | ?   | S<   | 08:08 | 0:00 | [kworker/0:0H] |
| root | 7   | 0.0  | 0.0  | 0     | 0    | ?   | S    | 08:08 | 0:05 | [rcu_sched]    |
| root | 8   | 0.0  | 0.0  | 0     | 0    | ?   | S    | 08:08 | 0:04 | [rcuos/0]      |
| root | 9   | 0.0  | 0.0  | 0     | 0    | ?   | S    | 08:08 | 0:01 | [rcuos/1]      |

- **USER** = user owning the process
- **PID** = process ID of the process
- **%CPU** = It is the CPU time used divided by the time the process has been running.
- **%MEM** = ratio of the process's resident set size to the physical memory on the machine
- **VSZ** = virtual memory usage of entire process (in KiB)
- **RSS** = resident set size, the non-swapped physical memory that a task has used (in KiB)
- **TTY** = controlling tty (terminal)
- **STAT** = multi-character process state
- **START** = starting time or date of the process
- **TIME** = cumulative CPU time
- **COMMAND** = command with all its arguments

```
ps -eF
```

-F option, which is the **extra full** format.

| UID  | PID | PPID | C | SZ   | RSS  | PSR | STIME | TTY | TIME     | CMD            |
|------|-----|------|---|------|------|-----|-------|-----|----------|----------------|
| root | 1   | 0    | 0 | 8477 | 3308 | 5   | 08:08 | ?   | 00:00:00 | /sbin/init     |
| root | 2   | 0    | 0 | 0    | 0    | 3   | 08:08 | ?   | 00:00:00 | [kthreadd]     |
| root | 3   | 2    | 0 | 0    | 0    | 0   | 08:08 | ?   | 00:00:00 | [ksoftirqd/0]  |
| root | 5   | 2    | 0 | 0    | 0    | 0   | 08:08 | ?   | 00:00:00 | [kworker/0:0H] |
| root | 7   | 2    | 0 | 0    | 0    | 3   | 08:08 | ?   | 00:00:05 | [rcu_sched]    |
| root | 8   | 2    | 0 | 0    | 0    | 0   | 08:08 | ?   | 00:00:04 | [rcuos/0]      |
| root | 9   | 2    | 0 | 0    | 0    | 0   | 08:08 | ?   | 00:00:01 | [rcuos/1]      |
| root | 10  | 2    | 0 | 0    | 0    | 0   | 08:08 | ?   | 00:00:01 | [rcuos/2]      |
| root | 11  | 2    | 0 | 0    | 0    | 1   | 08:08 | ?   | 00:00:01 | [rcuos/3]      |
| root | 12  | 2    | 0 | 0    | 0    | 2   | 08:08 | ?   | 00:00:01 | [rcuos/4]      |
| root | 13  | 2    | 0 | 0    | 0    | 0   | 08:08 | ?   | 00:00:01 | [rcuos/5]      |
| root | 14  | 2    | 0 | 0    | 0    | 0   | 08:08 | ?   | 00:00:00 | [rcuos/6]      |

-F option gives three additional columns of output.

**SZ** column gives the **virtual size of the process**.

**RSS** is the **resident set size**, the **non-swapped physical memory that the process has used** in kilobytes.

**PSR** indicates the **processor that the process is currently assigned to**.

```
ps -eLF
```

The -L option is for getting information about the threads

| UID  | PID | PPID | LWP | C | NLWP | SZ   | RSS  | PSR | STIME   | TTY | TIME     | CMD           |
|------|-----|------|-----|---|------|------|------|-----|---------|-----|----------|---------------|
| root | 1   | 0    | 1   | 0 | 1    | 830  | 1148 | 0   | Feb20   | ?   | 00:00:01 | /sbin/init    |
| root | 2   | 0    | 2   | 0 | 1    | 0    | 0    | 0   | Feb20   | ?   | 00:00:00 | [kthreadd]    |
| root | 3   | 2    | 3   | 0 | 1    | 0    | 0    | 0   | Feb20   | ?   | 00:02:23 | [ksoftirqd/0] |
| root | 6   | 2    | 6   | 0 | 1    | 0    | 0    | 0   | Feb20   | ?   | 00:00:00 | [migration/0] |
| root | 7   | 2    | 7   | 0 | 1    | 0    | 0    | 0   | 1 Feb20 | ?   | 00:00:00 | [migration/1] |
| root | 9   | 2    | 9   | 0 | 1    | 0    | 0    | 0   | 1 Feb20 | ?   | 00:13:50 | [ksoftirqd/1] |
| root | 11  | 2    | 11  | 0 | 1    | 0    | 0    | 0   | 0 Feb20 | ?   | 00:00:00 | [cpuset]      |
| root | 12  | 2    | 12  | 0 | 1    | 0    | 0    | 0   | 0 Feb20 | ?   | 00:00:00 | [khelper]     |
| root | 13  | 2    | 13  | 0 | 1    | 0    | 0    | 0   | 1 Feb20 | ?   | 00:00:00 | [netns]       |
| root | 14  | 2    | 14  | 0 | 1    | 0    | 0    | 0   | 1 Feb20 | ?   | 00:00:00 | [kworker/u:1] |
| .... |     |      |     |   |      |      |      |     |         |     |          |               |
| root | 842 | 1    | 842 | 0 | 1    | 1606 | 664  | 0   | Feb20   | ?   | 00:00:00 | /usr/sbin/ss  |

This gives one line of information per thread

LWP gives the **thread id of each thread**.

- NLWP is the **number of threads in the system** for the underlying process.
- The -L option gives an idea about the **total number of concurrent threads** in the system.

# **lsof**

lsof stands for List Open Files.

- It is easy to remember lsof command if you think of it as “**ls + of**”, where **ls** stands for list, and **of** stands for open files.
- list the information about the files that are opened by various processes.
- In linux everything is file , using lsof you can get the information about any opened files.
- lsof

| COMMAND | PID | USER | FD  | TYPE | DEVICE | SIZE/OFF | NODE NAME         |
|---------|-----|------|-----|------|--------|----------|-------------------|
| init    | 1   | root | cwd | DIR  | 8,1    | 4096     | 2 /               |
| init    | 1   | root | txt | REG  | 8,1    | 124704   | 917562 /sbin/init |
| init    | 1   | root | 0u  | CHR  | 1,3    | 0t0      | 4369 /dev/null    |
| init    | 1   | root | 1u  | CHR  | 1,3    | 0t0      | 4369 /dev/null    |
| init    | 1   | root | 2u  | CHR  | 1,3    | 0t0      | 4369 /dev/null    |
| init    | 1   | root | 3r  | FIFO | 0,8    | 0t0      | 6323 pipe         |
| ...     |     |      |     |      |        |          |                   |

- By default One file per line is displayed

FD – Represents the file descriptor. Some of the values of FDs are,

- cwd – Current Working Directory
- txt – Text file
- mem – Memory mapped file
- mmap – Memory mapped device
- NUMBER – Represent the actual file descriptor. The character after the number i.e ‘1u’, represents the mode in which the file is opened. r for read, w for write, u for read and write.

TYPE – Specifies the type of the file. Some of the values of TYPES are,

- REG – Regular File
- DIR – Directory
- FIFO – First In First Out
- CHR – Character special file

## List processes which opened a specific file

list only the processes which opened a specific file by providing the filename as arguments.

```
lsof /var/log/syslog
```

| COMMAND  | PID | USER   | FD | TYPE | DEVICE | SIZE/OFF | NODE   | NAME            |
|----------|-----|--------|----|------|--------|----------|--------|-----------------|
| rsyslogd | 488 | syslog | 1w | REG  | 8,1    | 1151     | 268940 | /var/log/syslog |

FD is followed by one of these characters, describing the mode under which the file is open:

- r for read access;
- w for write access;
- u for read and write access;

## List files opened by a specific user

```
lsof -u lakshmanan
```

| COMMAND   | PID  | USER       | FD  | TYPE | DEVICE | SIZE/OFF | NODE   | NAME             |
|-----------|------|------------|-----|------|--------|----------|--------|------------------|
| update-no | 1892 | lakshmanan | 20r | FIFO | 0,8    | 0t0      | 14536  | pipe             |
| update-no | 1892 | lakshmanan | 21w | FIFO | 0,8    | 0t0      | 14536  | pipe             |
| bash      | 1995 | lakshmanan | cwd | DIR  | 8,1    | 4096     | 393218 | /home/lakshmanan |

**list files opened by all users, expect some 1 or 2.**

```
lsof -u ^lakshmanan
```

| COMMAND   | PID  | USER  | FD  | TYPE | DEVICE | SIZE/OFF | NODE | NAME       |
|-----------|------|-------|-----|------|--------|----------|------|------------|
| rtkit-dae | 1380 | rtkit | 7u  | 0000 | 0,9    | 0        | 4360 | anon_inode |
| udisks-da | 1584 | root  | cwd | DIR  | 8,1    | 4096     | 2    | /          |

## List opened files under a directory

```
lsof +D /var/log/
```

| COMMAND   | PID | USER   | FD | TYPE | DEVICE | SIZE/OFF | NODE   | NAME                        |
|-----------|-----|--------|----|------|--------|----------|--------|-----------------------------|
| rsyslogd  | 488 | syslog | 1w | REG  | 8,1    | 1151     | 268940 | /var/log/syslog             |
| rsyslogd  | 488 | syslog | 2w | REG  | 8,1    | 2405     | 269616 | /var/log/auth.log           |
| console-k | 144 | root   | 9w | REG  | 8,1    | 10871    | 269369 | /var/log/ConsoleKit/history |

+D will recurse the sub directories also. If you don't want lsof to recurse, then use '+d' option.

## List opened files based on process names starting with

```
lsof -c ssh -c init
```

| COMMAND   | PID  | USER       | FD  | TYPE | DEVICE | SIZE/OFF | NODE    | NAME                              |
|-----------|------|------------|-----|------|--------|----------|---------|-----------------------------------|
| init      | 1    | root       | txt | REG  | 8,1    | 124704   | 917562  | /sbin/init                        |
| init      | 1    | root       | mem | REG  | 8,1    | 1434180  | 1442625 | /lib/i386-linux-gnu/libc-2.13.so  |
| init      | 1    | root       | mem | REG  | 8,1    | 30684    | 1442694 | /lib/i386-linux-gnu/librt-2.13.so |
| ...       |      |            |     |      |        |          |         |                                   |
| ssh-agent | 1528 | lakshmanan | 1u  | CHR  | 1,3    | 0t0      | 4369    | /dev/null                         |
| ssh-agent | 1528 | lakshmanan | 2u  | CHR  | 1,3    | 0t0      | 4369    | /dev/null                         |

## List all open files by a specific process

```
lsof -p 1753
```

**ss**

- used to show socket statistics
- It allows showing information similar to netstat command.
- The ss command gets its information directly from kernel space.
- The **ss command** is fast. Information is very fast when you are searching.
- gain more useful information about TCP and state information with the **ss command**.

**ss command** can provide information about:

- All TCP sockets.
- All UDP sockets.
- All established ssh / ftp / http / https connections.
- Filtering by state (such as connected, synchronized), addresses and ports.

## List all connection

ss | more

| Netid | State | Recv-Q | Send-Q | Local Address:Port           | Peer Address:Port   |
|-------|-------|--------|--------|------------------------------|---------------------|
| u_str | ESTAB | 0      | 0      | * 219827                     | * 220601            |
| u_str | ESTAB | 0      | 0      | @/tmp/.X11-unix/X0 223495    | * 223494            |
| u_str | ESTAB | 0      | 0      | * 220663                     | * 220044            |
| u_str | ESTAB | 0      | 0      | @/tmp/.X11-unix/X0 220453    | *219708             |
| u_str | ESTAB | 0      | 0      | @/tmp/dbus-eIWbPTCWIq 223497 | * 223496            |
| ..... |       |        |        |                              |                     |
| tcp   | ESTAB | 0      | 0      | 192.168.1.2:33643            | 68.232.35.139:https |
| tcp   | ESTAB | 0      | 0      | 192.168.1.2:56529            | 54.236.180.90:9999  |

- **Netid**, it denotes socket type and transport protocol
- **tcp, udp, raw, u\_str** is abbreviation for **unix\_stream**, **u\_dgr** for **UNIX datagram sockets**, **nl** for **netlink**, **p\_raw** and **p\_dgr** for **raw and datagram packet sockets**.
- **State**-Socket state ,the names are standard TCP names, except for UNCONN, which cannot happen for TCP, but normal for not connected sockets of another types.
- **Recv-Q and Send-Q** showing amount of data queued for receive and transmit.
- last two columns display **local address and port of the socket** and **its peer address**, if the socket is connected.

## List all TCP connection

To list all TCP connection use -t option:

**ss -t**

| State | Recv-Q | Send-Q | Local Address:Port | Peer Address:Port    |
|-------|--------|--------|--------------------|----------------------|
| ESTAB | 0      | 0      | 172.17.7.47:38272  | 172.17.5.2:ssh       |
| ESTAB | 0      | 0      | 172.17.7.47:57092  | 74.125.204.189:https |
| ESTAB | 0      | 0      | 172.17.7.47:38091  | 204.11.109.67:http   |
| ESTAB | 0      | 0      | 172.17.7.47:53200  | 74.125.130.109:imaps |

“-t” options report only “ESTABLISHED” or “CONNECTED” connections. If use the “-ta” option report show all TCP connections (connected and listening).

**ss -ta**

| State  | Recv-Q | Send-Q | Local Address:Port | Peer Address:Port   |
|--------|--------|--------|--------------------|---------------------|
| LISTEN | 0      | 128    | *:ssh              | *:*                 |
| LISTEN | 0      | 128    | 127.0.0.1:ipp      | *:*                 |
| LISTEN | 0      | 50     | 127.0.0.1:mysql    | *:*                 |
| ESTAB  | 0      | 0      | 172.17.7.47:53700  | 216.58.197.34:https |
| ESTAB  | 0      | 0      | 172.17.7.47:38272  | 172.17.5.2:ssh      |

## List all UDP connection

```
ss -ua
```

| State | Recv-Q | Send-Q | Local Address:Port | Peer Address:Port |
|-------|--------|--------|--------------------|-------------------|
|-------|--------|--------|--------------------|-------------------|

|        |   |   |       |     |
|--------|---|---|-------|-----|
| UNCONN | 0 | 0 | *:ipp | *:* |
|--------|---|---|-------|-----|

|        |   |   |        |     |
|--------|---|---|--------|-----|
| UNCONN | 0 | 0 | *:mdns | *:* |
|--------|---|---|--------|-----|

|        |   |   |        |     |
|--------|---|---|--------|-----|
| UNCONN | 0 | 0 | *:mdns | *:* |
|--------|---|---|--------|-----|

|        |   |   |         |     |
|--------|---|---|---------|-----|
| UNCONN | 0 | 0 | *:35057 | *:* |
|--------|---|---|---------|-----|

|        |   |   |         |      |
|--------|---|---|---------|------|
| UNCONN | 0 | 0 | ::46110 | ...* |
|--------|---|---|---------|------|

|        |   |   |        |      |
|--------|---|---|--------|------|
| UNCONN | 0 | 0 | ::mdns | ...* |
|--------|---|---|--------|------|

## Do not resolve hostname

- use the "n" option to prevent ss from resolving ip addresses to hostnames.

### ss -ua

| • | State  | Recv-Q | Send-Q | Local Address:Port | Peer Address:Port |
|---|--------|--------|--------|--------------------|-------------------|
| • | UNCONN | 0      | 0      | *:ipp              | .*                |
| • | UNCONN | 0      | 0      | *:mdns             | .*                |
| • | UNCONN | 0      | 0      | *:mdns             | .*                |
| • | UNCONN | 0      | 0      | *:35057            | .*                |
| • | UNCONN | 0      | 0      | ::46110            | ...*              |
| • | UNCONN | 0      | 0      | ::mdns             | ...*              |

### ss -nua

| • | State  | Recv-Q | Send-Q | Local Address:Port | Peer Address:Port |
|---|--------|--------|--------|--------------------|-------------------|
| • | UNCONN | 0      | 0      | *:631              | .*                |
| • | UNCONN | 0      | 0      | *:5353             | .*                |
| • | UNCONN | 0      | 0      | *:5353             | .*                |
| • | UNCONN | 0      | 0      | *:35057            | .*                |
| • | UNCONN | 0      | 0      | ::46110            | ...*              |
| • | UNCONN | 0      | 0      | ::5353             | ...*              |

## Show only listening sockets

list out all the listening sockets. For example apache web server opens a socket connection on port 80 to listen for incoming connections.

### ss -ltn (listening tcp connections)

| • | State  | Recv-Q | Send-Q | Local Address:Port | Peer Address:Port |
|---|--------|--------|--------|--------------------|-------------------|
| • | LISTEN | 0      | 128    | *:22               | .*.*              |
| • | LISTEN | 0      | 128    | 127.0.0.1:631      | .*.*              |
| • | LISTEN | 0      | 50     | 127.0.0.1:3306     | .*.*              |
| • | LISTEN | 0      | 128    | :::22              | ...:*             |
| • | LISTEN | 0      | 128    | ::1:631            | ...:*             |

### ss -lun (listening udp connections)

| • | State  | Recv-Q | Send-Q | Local Address:Port | Peer Address:Port |
|---|--------|--------|--------|--------------------|-------------------|
| • | UNCONN | 0      | 0      | *:631              | .*.*              |
| • | UNCONN | 0      | 0      | *:5353             | .*.*              |
| • | UNCONN | 0      | 0      | *:5353             | .*.*              |
| • | UNCONN | 0      | 0      | *:35057            | .*.*              |
| • | UNCONN | 0      | 0      | :::46110           | ...:*             |
| • | UNCONN | 0      | 0      | :::5353            | ...:*             |

## Print process name and pid

To print out the process name/pid which owns the connection use the p option

### ss -ltp

| State  | Recv-Q | Send-Q | Local Address:Port | Peer Address:Port                                   |
|--------|--------|--------|--------------------|-----------------------------------------------------|
| LISTEN | 0      | 100    | 127.0.0.1:smtp     | *:*                                                 |
| LISTEN | 0      | 128    | 127.0.0.1:9050     | *:*                                                 |
| LISTEN | 0      | 128    | *:90               | *:*                                                 |
| LISTEN | 0      | 128    | *:db-lsp           | *:* users:(("dropbox",3566,32))                     |
| LISTEN | 0      | 5      | 127.0.0.1:6600     | *:*                                                 |
| LISTEN | 0      | 128    | 127.0.0.1:9000     | *:* users:(("php5-fpm",1620,0),("php5-fpm",1619,0)) |

## Print summary statistics

The s option prints out the statistics.

### ss -s

Total: 913 (kernel 0)

TCP: 31 (estab 24, closed 0, orphaned 0, synrecv 0, timewait 0/0), ports 0

Transport Total IP IPv6

|      |    |    |   |
|------|----|----|---|
| *    | 0  | -  | - |
| RAW  | 0  | 0  | 0 |
| UDP  | 6  | 4  | 2 |
| TCP  | 31 | 28 | 3 |
| INET | 37 | 32 | 5 |
| FRAG | 0  | 0  | 0 |

## Display timer information

With the '-o' option, the time information of each connection would be displayed.

**\$ ss -tn -o**

| State | Recv-Q | Send-Q | Local Address:Port | Peer Address:Port                           |
|-------|--------|--------|--------------------|---------------------------------------------|
| ESTAB | 0      | 0      | 192.168.1.2:43839  | 108.160.162.37:80                           |
| ESTAB | 0      | 0      | 192.168.1.2:36335  | 204.144.140.26:80 timer:(keepalive,26sec,0) |
| ESTAB | 0      | 0      | 192.168.1.2:33141  | 83.170.73.249:6667                          |
| ESTAB | 0      | 0      | 192.168.1.2:58857  | 74.121.141.84:80 timer:(keepalive,23sec,0)  |
| ESTAB | 0      | 0      | 192.168.1.2:42794  | 173.194.40.239:80 timer:(keepalive,32sec,0) |

## Display only IPv4 or IPv6 socket connections

**ss -tl4**

- State Recv-Q Send-Q Local Address:Port Peer Address:Port
- LISTEN 0 128 \*:ssh \*:\*
- LISTEN 0 128 127.0.0.1:ipp \*:\*
- LISTEN 0 50 127.0.0.1:mysql \*:\*

**ss -tl6**

- State Recv-Q Send-Q Local Address:Port Peer Address:Port
- LISTEN 0 128 ::ssh ::\*
- LISTEN 0 128 ::1:ipp ::\*

## Filtering connections by tcp state

ss command supports filters that can be used to display only specific connections.

### Display sockets with state time-wait

**ss -t4 state time-wait**

| Recv-Q | Send-Q | Local Address:Port | Peer Address:Port   |
|--------|--------|--------------------|---------------------|
| 0      | 0      | 192.168.1.2:42261  | 199.59.150.39:https |
| 0      | 0      | 127.0.0.1:43541    | 127.0.0.1:2633      |

### Filter connections by address and port number

Filter by destination port

**ss -nt dst :443 or dst :80**

Filter by destination address

**ss -nt dst 74.125.236.178**

Address and Port combined

**\$ ss -nt dst 74.125.236.178:80**

# **laptop**

used to monitor the Disk I/O

- iotop is an open source and free utility
- used to monitor the Disk I/O and to trace the exact process or high used user disk read/writes.
- iotop tool is based on Python programming
- Need sudo privilege to run iotop

## Install iotop

### Centos ,redhat

```
yum install iotop
```

### **Debian ,ubuntu**

```
sudo apt-get install iotop
```

root@txvip1:~

File Edit View Terminal Help


Total DISK READ: 0.00 B/s | Total DISK WRITE: 149.47 K/s

| TID   | PRIo | USER  | DISK READ | DISK WRITE | SWAPIN | IO>    | COMMAND                                  |
|-------|------|-------|-----------|------------|--------|--------|------------------------------------------|
| 417   | be/3 | root  | 0.00 B/s  | 0.00 B/s   | 0.00 % | 0.03 % | [jbd2/sda5-8]                            |
| 42451 | be/4 | mysql | 0.00 B/s  | 26.83 K/s  | 0.00 % | 0.01 % | mysqld --basedir=/u-lib/mysql/mysql.sock |
| 27792 | be/4 | mysql | 0.00 B/s  | 23.00 K/s  | 0.00 % | 0.01 % | mysqld --basedir=/u-lib/mysql/mysql.sock |
| 43128 | be/4 |       | 0.00 B/s  | 7.67 K/s   | 0.00 % | 0.00 % | lighttpd -f /etc/lighttpd.conf           |
| 19380 | be/4 |       | 0.00 B/s  | 5.58 M/s   | 0.00 % | 0.00 % | php-cgi                                  |
| 42395 | be/4 | mysql | 0.00 B/s  | 0.00 B/s   | 0.00 % | 0.00 % | mysqld --basedir=/u-lib/mysql/mysql.sock |
| 1     | be/4 | root  | 0.00 B/s  | 0.00 B/s   | 0.00 % | 0.00 % | init                                     |
| 2     | be/4 | root  | 0.00 B/s  | 0.00 B/s   | 0.00 % | 0.00 % | [kthreadd]                               |
| 3     | rt/4 | root  | 0.00 B/s  | 0.00 B/s   | 0.00 % | 0.00 % | [migration/0]                            |
| 4     | be/4 | root  | 0.00 B/s  | 0.00 B/s   | 0.00 % | 0.00 % | [ksoftirqd/0]                            |
| 5     | rt/4 | root  | 0.00 B/s  | 0.00 B/s   | 0.00 % | 0.00 % | [migration/0]                            |
| 6     | rt/4 | root  | 0.00 B/s  | 0.00 B/s   | 0.00 % | 0.00 % | [watchdog/0]                             |
| 7     | rt/4 | root  | 0.00 B/s  | 0.00 B/s   | 0.00 % | 0.00 % | [migration/1]                            |
| 8     | rt/4 | root  | 0.00 B/s  | 0.00 B/s   | 0.00 % | 0.00 % | [migration/1]                            |
| 9     | be/4 | root  | 0.00 B/s  | 0.00 B/s   | 0.00 % | 0.00 % | [ksoftirqd/1]                            |
| 10    | rt/4 | root  | 0.00 B/s  | 0.00 B/s   | 0.00 % | 0.00 % | [watchdog/1]                             |
| 11    | rt/4 | root  | 0.00 B/s  | 0.00 B/s   | 0.00 % | 0.00 % | [migration/2]                            |
| 12    | rt/4 | root  | 0.00 B/s  | 0.00 B/s   | 0.00 % | 0.00 % | [migration/2]                            |
| 13    | be/4 | root  | 0.00 B/s  | 0.00 B/s   | 0.00 % | 0.00 % | [ksoftirqd/2]                            |
| 14    | rt/4 | root  | 0.00 B/s  | 0.00 B/s   | 0.00 % | 0.00 % | [watchdog/2]                             |
| 15    | rt/4 | root  | 0.00 B/s  | 0.00 B/s   | 0.00 % | 0.00 % | [migration/3]                            |
| 16    | rt/4 | root  | 0.00 B/s  | 0.00 B/s   | 0.00 % | 0.00 % | [migration/3]                            |
| 17    | be/4 | root  | 0.00 B/s  | 0.00 B/s   | 0.00 % | 0.00 % | [ksoftirqd/3]                            |
| 18    | rt/4 | root  | 0.00 B/s  | 0.00 B/s   | 0.00 % | 0.00 % | [watchdog/3]                             |
| 19    | rt/4 | root  | 0.00 B/s  | 0.00 B/s   | 0.00 % | 0.00 % | [migration/4]                            |
| 20    | rt/4 | root  | 0.00 B/s  | 0.00 B/s   | 0.00 % | 0.00 % | [migration/4]                            |
| 21    | be/4 | root  | 0.00 B/s  | 0.00 B/s   | 0.00 % | 0.00 % | [ksoftirqd/4]                            |
| 22    | rt/4 | root  | 0.00 B/s  | 0.00 B/s   | 0.00 % | 0.00 % | [watchdog/4]                             |

- The "IO" column lists total I/O for each process
- The "SwapIn" column lists **swap activity for each process.**
- TID stands for Thread ID

## iotop --only

to see only **processes or threads actually doing I/O**, instead of showing all processes or threads



|       |      | Total DISK READ: 0.00 B/s   Total DISK WRITE: 253.27 K/s |           |             |        |        |                                          |
|-------|------|----------------------------------------------------------|-----------|-------------|--------|--------|------------------------------------------|
| TID   | PRIo | USER                                                     | DISK READ | DISK WRITE  | SWAPIN | IO>    | COMMAND                                  |
| 417   | be/3 | root                                                     | 0.00 B/s  | 3.84 K/s    | 0.00 % | 0.03 % | [jbd2/sda5-8]                            |
| 42467 | be/4 | mysql                                                    | 0.00 B/s  | 65.24 K/s   | 0.00 % | 0.01 % | mysqld --basedir=/u/lib/mysql/mysql.sock |
| 27794 | be/4 | mysql                                                    | 0.00 B/s  | 30.70 K/s   | 0.00 % | 0.01 % | mysqld --basedir=/u/lib/mysql/mysql.sock |
| 14041 | be/4 | mysql                                                    | 0.00 B/s  | 26.86 K/s   | 0.00 % | 0.01 % | mysqld --basedir=/u/lib/mysql/mysql.sock |
| 19747 | be/4 |                                                          | 0.00 B/s  | 1868.84 K/s | 0.00 % | 0.00 % | php-cgi                                  |

1. Hit the **left** and **right** arrow keys to change the sorting.
2. Hit **r** to reverse the sorting order.
3. Hit **o** only to see processes or threads actually doing I/O, instead of showing all processes or threads.
4. Hit **p** only show processes. Normally iotop shows all threads.

**slabtop**

- slab - **memory management mechanism** intended for the **efficient memory allocation of kernel objects**
- slabtop command displays detailed **kernel slab cache information** in real time

```
• Active / Total Objects (% used) : 216141 / 216625 (99.8%)
 Active / Total Slabs (% used) : 6841 / 6841 (100.0%)
 Active / Total Caches (% used) : 67 / 104 (64.4%)
 Active / Total Size (% used) : 35774.17K / 36179.05K (98.9%)
 Minimum / Average / Maximum Object : 0.01K / 0.17K / 8.00K
```

| OBJS  | ACTIVE | USE  | OBJ   | SIZE | SLABS | OBJ/SLAB | CACHE             | SIZE | NAME |
|-------|--------|------|-------|------|-------|----------|-------------------|------|------|
| 46016 | 46016  | 100% | 0.12K | 1438 | 32    | 5752K    | dentry            |      |      |
| 37887 | 37887  | 100% | 0.05K | 519  | 73    | 2076K    | buffer_head       |      |      |
| 23426 | 23426  | 100% | 0.59K | 1802 | 13    | 14416K   | ext4_inode_cache  |      |      |
| 17024 | 16867  | 99%  | 0.03K | 133  | 128   | 532K     | kmalloc-32        |      |      |
| 16320 | 16320  | 100% | 0.06K | 255  | 64    | 1020K    | kmalloc-64        |      |      |
| 14658 | 14658  | 100% | 0.09K | 349  | 42    | 1396K    | kmalloc-96        |      |      |
| 11264 | 11264  | 100% | 0.01K | 22   | 512   | 88K      | kmalloc-8         |      |      |
| 7740  | 7740   | 100% | 0.33K | 645  | 12    | 2580K    | inode_cache       |      |      |
| 7424  | 7424   | 100% | 0.03K | 58   | 128   | 232K     | anon_vma          |      |      |
| 4788  | 4678   | 97%  | 0.19K | 228  | 21    | 912K     | kmalloc-192       |      |      |
| 4303  | 4303   | 100% | 0.30K | 331  | 13    | 1324K    | radix_tree_node   |      |      |
| 4059  | 4045   | 99%  | 0.36K | 369  | 11    | 1476K    | proc_inode_cache  |      |      |
| 3840  | 3840   | 100% | 0.02K | 15   | 256   | 60K      | kmalloc-16        |      |      |
| 3584  | 3584   | 100% | 0.01K | 7    | 512   | 28K      | ext4_io_page      |      |      |
| 1615  | 1615   | 100% | 0.05K | 19   | 85    | 76K      | Acpi-State        |      |      |
| 1360  | 1360   | 100% | 0.02K | 8    | 170   | 32K      | nsproxy           |      |      |
| 1216  | 1216   | 100% | 0.12K | 38   | 32    | 152K     | kmalloc-128       |      |      |
| 1088  | 1088   | 100% | 0.06K | 17   | 64    | 68K      | journal_head      |      |      |
| 1000  | 970    | 97%  | 0.50K | 125  | 8     | 500K     | kmalloc-512       |      |      |
| 957   | 957    | 100% | 0.36K | 87   | 11    | 348K     | shmem_inode_cache |      |      |
| 850   | 850    | 100% | 0.02K | 5    | 170   | 20K      | extent_status     |      |      |
| 676   | 676    | 100% | 0.15K | 26   | 26    | 104K     | idr_layer_cache   |      |      |
| 672   | 608    | 90%  | 0.25K | 42   | 16    | 168K     | kmalloc-256       |      |      |
| 420   | 420    | 100% | 0.38K | 42   | 10    | 168K     | sock_inode_cache  |      |      |
| 357   | 357    | 100% | 0.08K | 7    | 51    | 28K      | task_delay_info   |      |      |

**vmstat (*virtual memory statistics*)**

- vmstat is a tool in Unix/Linux which is used to **Report virtual memory statistics.**
- It shows how much virtual memory there is
- how much is free and paging activity.
- page-ins and page-outs as they happen.

vmstat

procs -----memory----- --swap-- -----io---- -system-- -----cpu-----

| r | b | swpd | free   | buff   | cache   | si | so | bi | bo | in | cs  | us | sy | id | wa | st |
|---|---|------|--------|--------|---------|----|----|----|----|----|-----|----|----|----|----|----|
| 2 | 0 | 0    | 509288 | 218380 | 3525292 | 0  | 0  | 17 | 48 | 60 | 259 | 2  | 0  | 98 | 1  | 0  |

- **r:** The number of processes waiting for run time
- **b:** The number of processes in uninterruptible sleep. (b=blocked queue, waiting for resource)
- **swpd:** shows how **many blocks are swapped out to disk** (paged). The **amount of Virtual memory used.**
- **free:** The amount of Idle Memory
- **buff:** Memory used as buffers, like before/after I/O operations
- **cache:** Memory used as cache by the Operating System

**si:** Amount of memory swapped in from disk (/s). This shows page-ins

**so:** Amount of memory swapped to disk (/s). This shows page-outs

**Under IO we have:**

bi: Blocks received from block device - Read

bo: Blocks sent to a block device – Write

**Under System we have:**

in: The number of interrupts per second,

cs: The number of context switches per second

## Under CPU we have:

These are **percentages of total CPU time**.

us: % of CPU time spent in user mode

sy: % of CPU time spent running kernel code. (system time)

id: % of CPU idle time

wa: % of CPU time spent waiting for IO.

## /proc/meminfo

**Active** — The total amount of buffer or page cache memory, in kilobytes, that is in active use. This is memory that has been recently used and is usually **not reclaimed** for other purposes.

**Inactive** — The total amount of buffer or page cache memory, in kilobytes, that are free and available. This is memory that has not been recently used and can be **reclaimed** for other purposes.

**Dirty** — The total amount of memory, in kilobytes, waiting to be written back to the disk.

**Slab** — The total amount of memory, in kilobytes, used by the kernel to cache data structures for its own use.

- **memory management mechanism** intended for the **efficient memory allocation of kernel objects**

**SwapCached** — The amount of swap, in kilobytes, used as cache memory.

**Buffer** - data buffer (or just buffer) is a region of a physical memory storage used to temporarily store data while it is being moved from one place to another.

The data is stored in a buffer as it is retrieved from an input device (such as a microphone) or just before it is sent to an output device (such as speakers).

**Cache**-to reduce the average time to access data from the main memory. The cache is a smaller, faster memory which stores copies of the data from frequently used main memory locations. Most CPUs have different independent caches, including instruction and data caches, where the data cache is usually organized as a hierarchy of more cache levels (L1, L2, etc.).

When the processor needs to read from or write to a location in main memory, it first checks whether a copy of that data is in the cache. If so, the processor immediately reads from or writes to the cache, which is much faster than reading from or writing to main memory.

instruction cache to speed up executable instruction fetch

data cache to speed up data fetch and store,

translation lookaside buffer (TLB) used to speed up virtual-to-physical address translation for both executable instructions and data

# /proc/partitions

This **file** contains **partition block allocation information**.

Block devices are disk devices for which the kernel provides caching.

| major | minor | #blocks  | name |
|-------|-------|----------|------|
| 3     | 0     | 19531250 | hda  |
| 3     | 1     | 104391   | hda1 |
| 3     | 2     | 19422585 | hda2 |
| 253   | 0     | 22708224 | dm-0 |
| 253   | 1     | 524288   | dm-1 |

major – **The major number of the device with this partition.** The major number in the /proc/partitions, (3), corresponds with the block device ide0, in /proc/devices.

minor – **The minor number of the device with this partition.** This serves to separate the partitions into different physical devices and relates to the number at the end of the name of the partition.

#blocks – **Lists the number of physical disk blocks contained in a particular partition.**

name – **The name of the partition.**

# /proc/devices

This file **displays the various character and block devices currently configured**

Character devices:

1 mem

4 /dev/vc/0

4 tty

4 ttyS

5 /dev/tty

5 /dev/console

5 /dev/ptmx

7 vcs

10 misc

13 input

29 fb

36 netlink

128 ptm

136 pts

180 usb

Block devices:

1 ramdisk

3 ide0

9 md

22 ide1

253 device-mapper

254 mdp

The **output from /proc/devices includes the major number and name of the device**, and is broken into two major sections: **Character devices and Block devices**.

A Character ('c') Device is one with which the **Driver communicates by sending and receiving single characters** (bytes, octets). A **Block ('b')Device** is one with which the **Driver communicates by sending entire blocks of data**.

# Clearing swap space

free

```
total used free shared buffers cached
Mem: 7987492 7298164 689328 0 30416 457936
-/+ buffers/cache: 6809812 1177680
Swap: 5963772 609452 5354320
```

Use swapoff -a command

free

```
total used free shared buffers cached
Mem: 7987492 7777912 209580 0 39332 489864
-/+ buffers/cache: 7248716 738776
Swap: 0 0 0
```

And to re-enable it:

```
$ swapon -a
```

## Context switch

A **context switch** occurs when a computer's CPU switches from one process or thread to a different process or thread. Context switching allows for one CPU to handle numerous processes or threads without the need for additional processors.

## Interrupt

**interrupt** is a signal to the processor emitted by hardware or software indicating an event that needs immediate attention.


An interrupt alerts the processor to a high-priority condition requiring the interruption of the current code the processor is executing.

**Hardware interrupts** are used by devices to communicate that they require attention from the operating system.

A **software interrupt** is caused either by an exceptional condition in the processor itself, or a special instruction in the instruction set which

# Thread

*thread* of execution is the smallest sequence of programmed instructions that can be managed independently, which is typically a part of the operating system.



A process with two threads of execution, running on one processor.

## Clearing cache in linux

To free pagecache:

```
echo 1 > /proc/sys/vm/drop_caches
```

To free dentries and inodes:

```
echo 2 > /proc/sys/vm/drop_caches
```

To free pagecache, dentries and inodes:

```
echo 3 > /proc/sys/vm/drop_caches
```

An **inode** in your context is a **data structure that represents a file**. A **dentries** is a **data structure that represents a directory**. These structures could be **used to build a memory cache that represents the file structure on a disk**. To get a directory listing, the OS could go to the dentries--if the directory is there--list its contents (a series of inodes). If not there, go to the disk and read it into memory so that it can be used again.

---

**Page Cache** accelerates many accesses to files.

This happens because, when it first reads from or writes to data media like hard drives, Linux also stores data in **unused areas of memory**, which acts as a cache. If this data is read again later, it can be quickly read from this cache in memory

## Writing

If data is written, it is first written to the Page Cache and managed as one of its ***dirty pages***. *Dirty* means that the data is stored in the Page Cache, but needs to be written to the underlying storage device first. The content of these ***dirty pages*** is periodically transferred to the underlying storage device.

## Reading

**File blocks are written to the Page Cache not just during writing, but also when reading files.** For example, when you read a 100-megabyte file twice, once after the other, the second access will be quicker, because the file blocks come directly from the Page Cache in memory and do not have to be read from the hard disk again.